



Applied Natural Language Processing


Info 256

Lecture 20: Sequence labeling (April 9, 2019)


David Bamman, UC Berkeley

POS tagging

Labeling the tag that's correct
for the context.



Fruit flies like a banana



Time flies like an arrow

(Just tags in evidence within the Penn Treebank — more are possible!)

Named entity recognition

B-PERS I-PERS O O O O ORG

tim cook is the ceo of apple

3 or 4-class:

- person
- location
- organization
- (misc)

7-class:

- person
- location
- organization
- time
- money
- percent
- date

Supersense tagging

☐ B-artifact ☐ I-artifact ☒ B-motion ☐ B-time ☐ ☐ ☐ ☐ B-group

The station wagons arrived at noon, a long shining line

☐ B-motion ☐ ☐ B-location ☐ I-location

that coursed through the west campus.

1	person	7	cognition	13	attribute	19	quantity	25	plant
2	communication	8	possession	14	object	20	motive	26	relation
3	artifact	9	location	15	process	21	animal		
4	act	10	substance	16	Tops	22	body		
5	group	11	state	17	phenomenon	23	feeling		
6	food	12	time	18	event	24	shape		

Segmentation

- B = character is the start of new word
- I = character is inside existing word

#	b	l	a	c	k	l	i	v	e	s	m	a	t	t	e	r
B	B	I	I	I	I	B	I	I	I	I	B	I	I	I	I	I
B	B	I	I	I	I	B	I	I	I	B	I	I	I	I	I	I

black lives matter

black live smatter

Sequence labeling

$$x = \{x_1, \dots, x_n\}$$

$$y = \{y_1, \dots, y_n\}$$


- For a set of inputs x with n sequential time steps, one corresponding label y_i for each x_i


Sequence labeling models

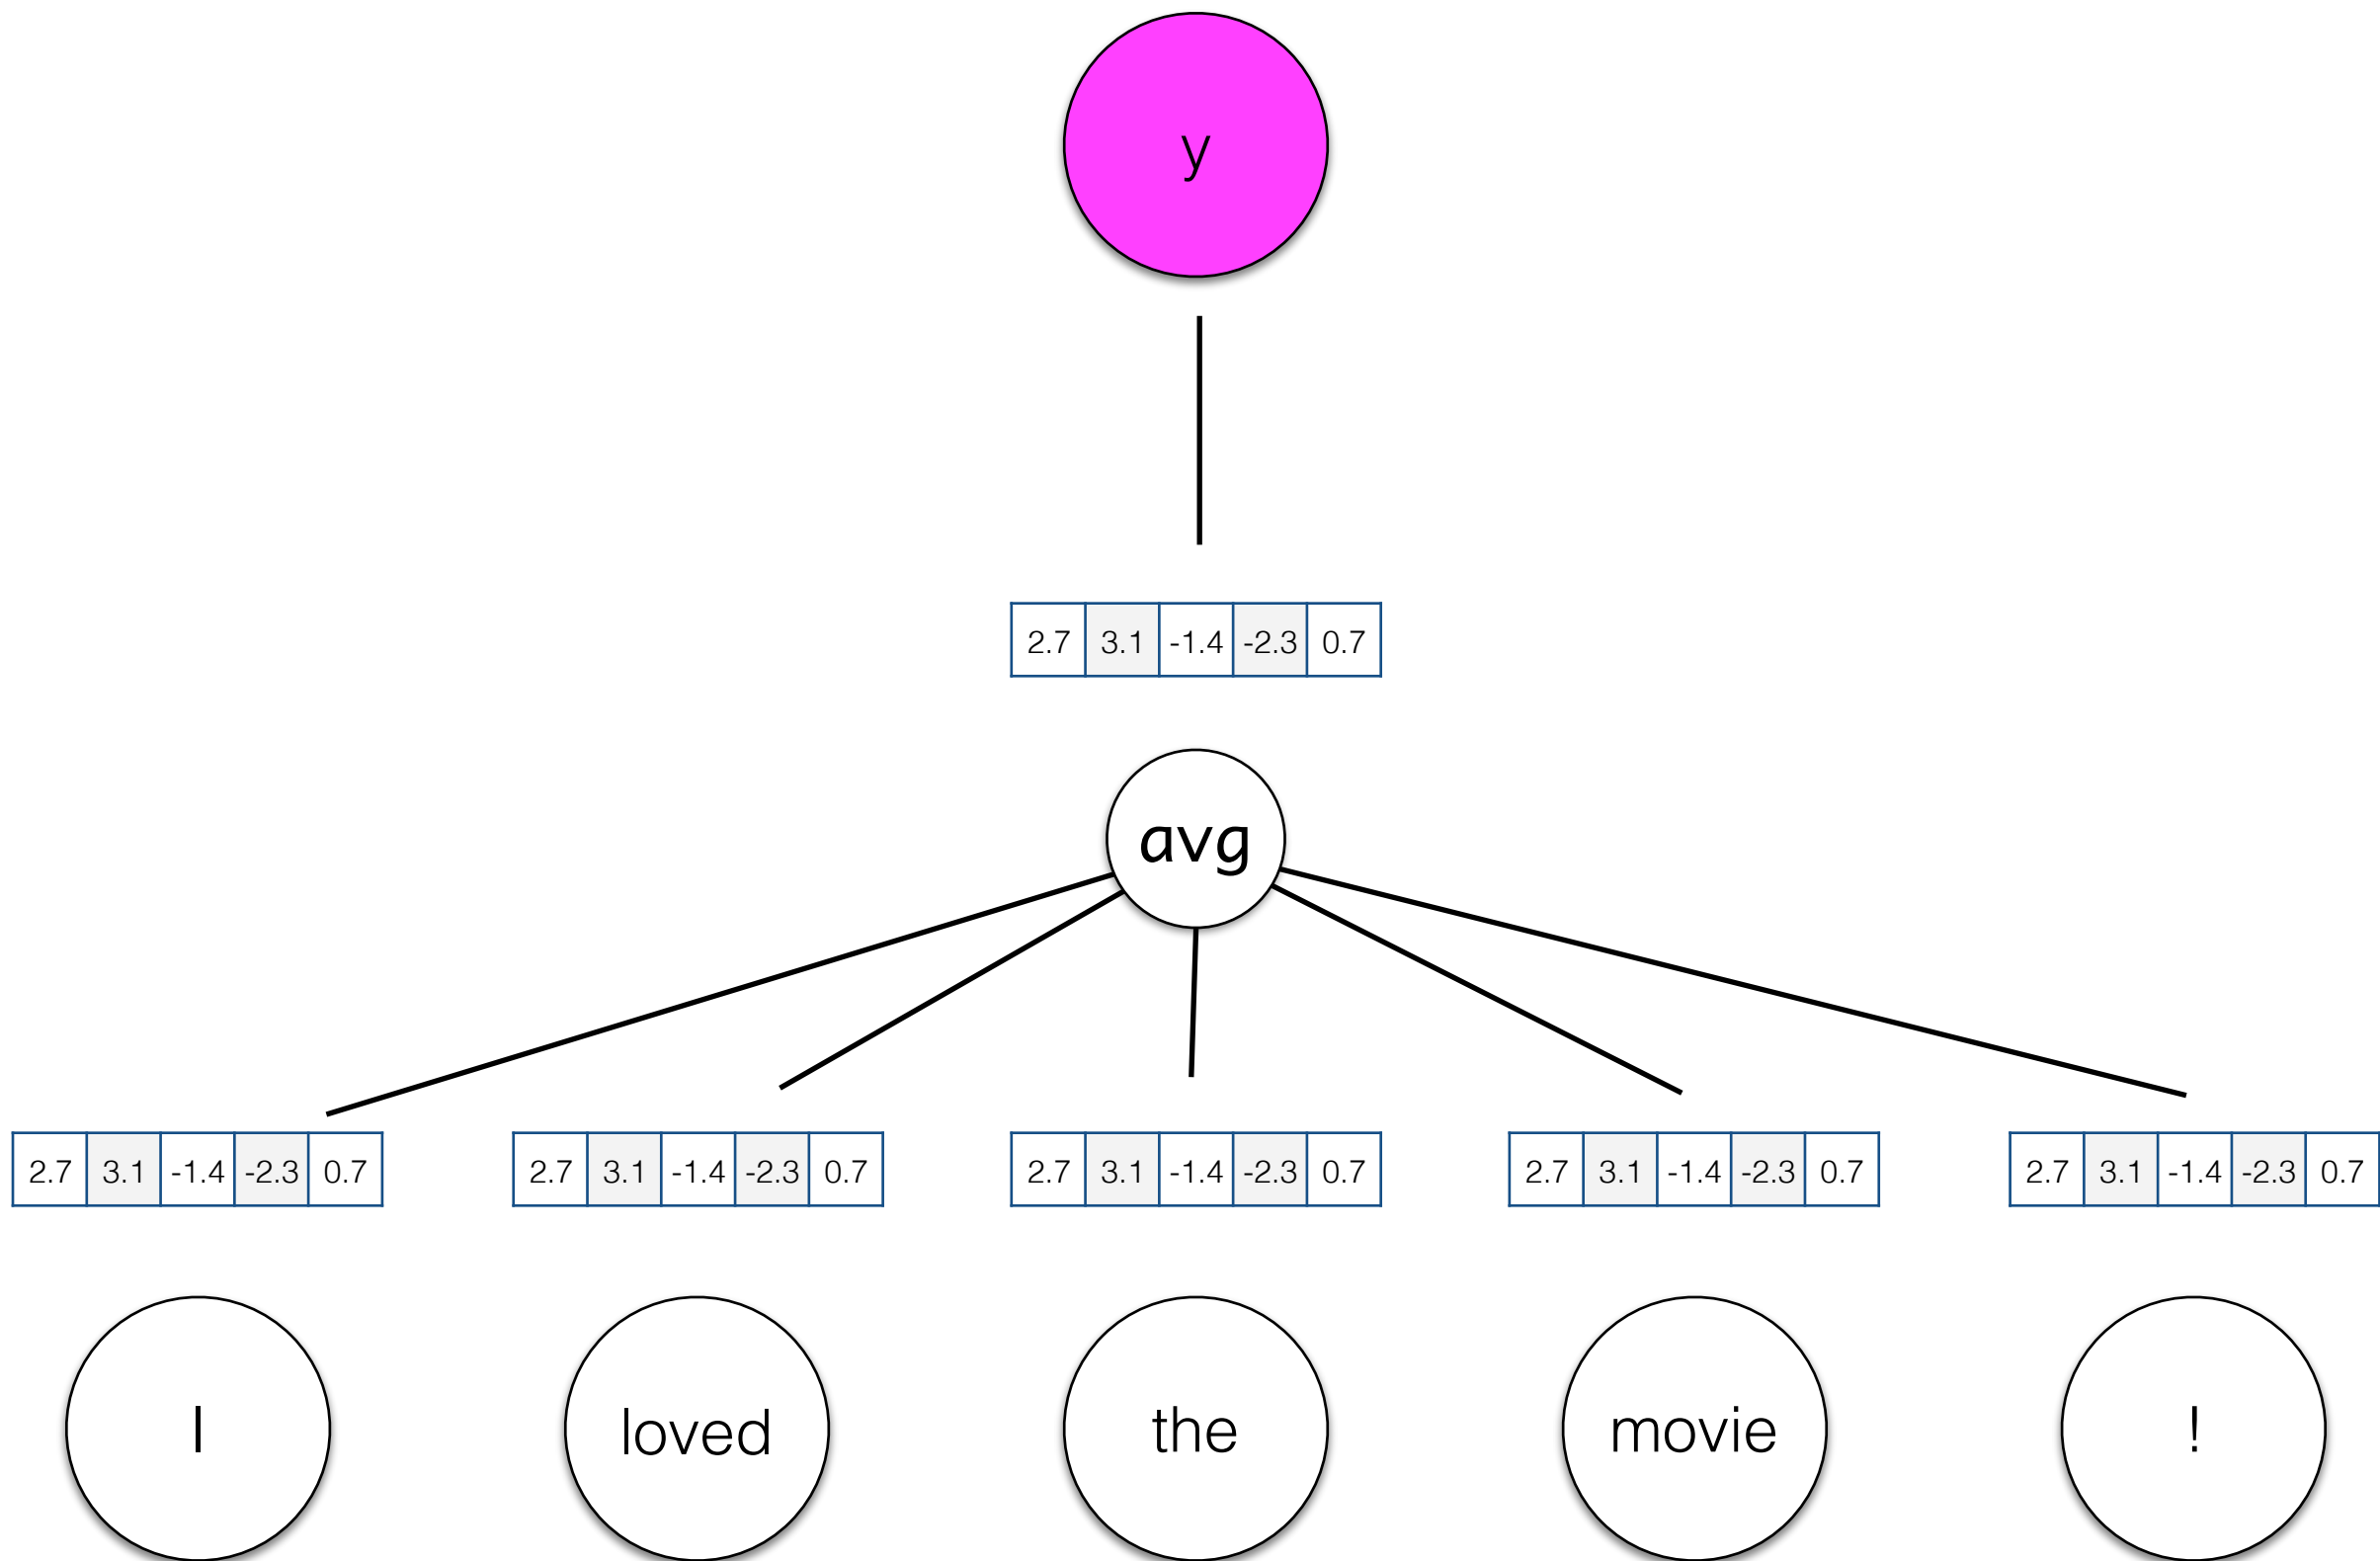
model	form	label dependency	rich features?
Hidden Markov Models	$\prod_{i=1}^N P(x_i y_i) P(y_i y_{i-1})$	Markov assumption	no
MEMM	$\prod_{i=1}^N P(y_i y_{i-1}, x, \beta)$	Markov assumption	yes
CRF	$P(y x, \beta)$	pairwise through entire sequence	yes
RNN	$\prod_{i=1}^N P(y_i x_{1:i}, \beta)$	none	distributed


Back to RNNs

- RNN allow arbitrarily-sized conditioning contexts; condition on the **entire sequence history**.
- We used RNNs for document classification to generate a **representation** of a sequence that we can then use for prediction.



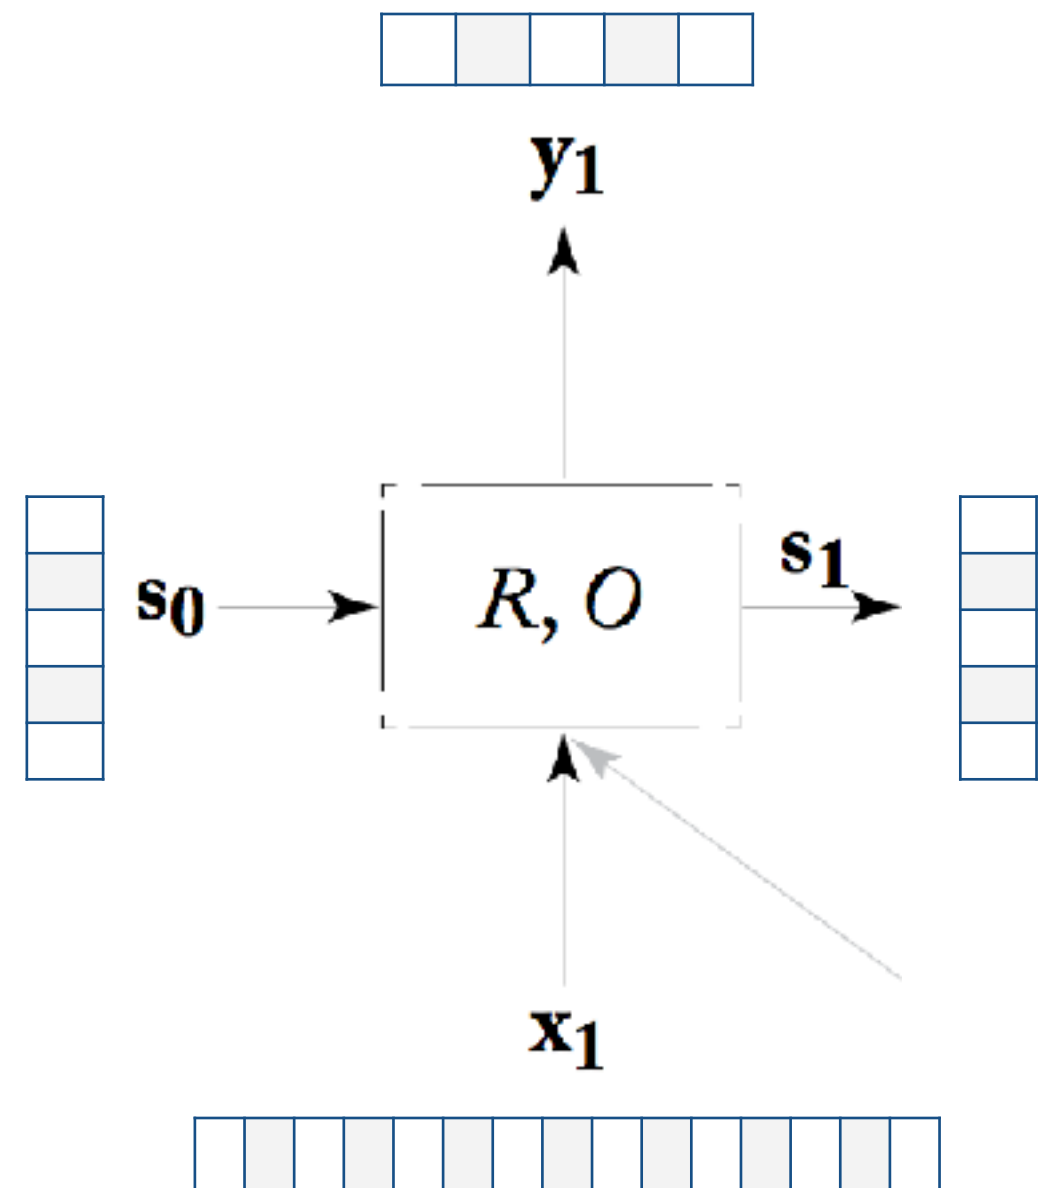






Recurrent neural network

- Each time step has two inputs:
 - x_i (the observation at time step i); one-hot vector, feature vector or **word embedding**.
 - s_{i-1} (the output of the previous state); base case: $s_0 = 0$ vector



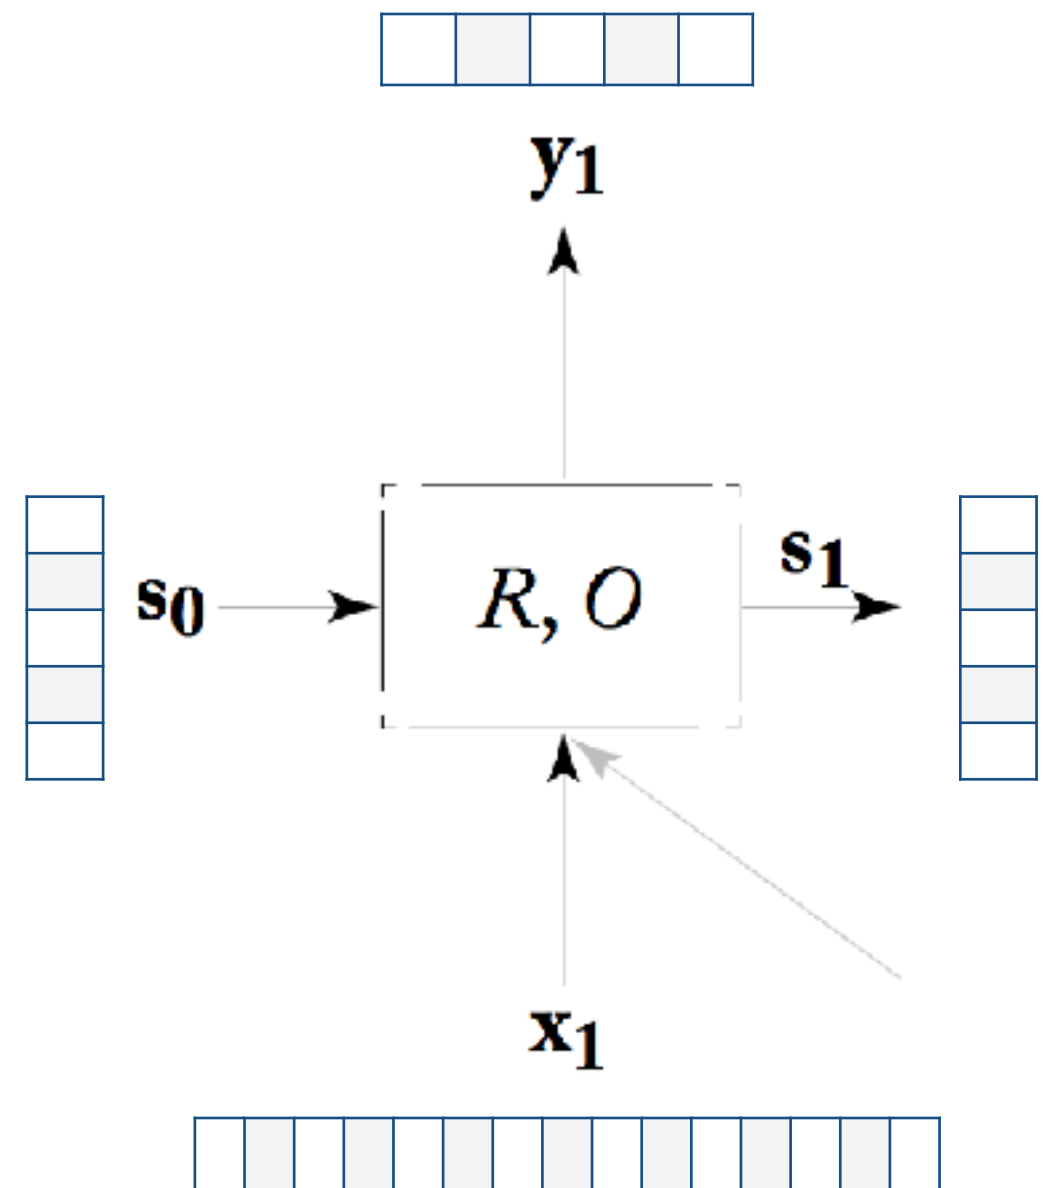
Recurrent neural network

$$s_i = R(x_i, s_{i-1})$$

R computes the output state as a function of the current input and previous state

$$y_i = O(s_i)$$

O computes the output as a function of the current output state



“Simple” RNN

$g = \tanh$ or relu

$$s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$$


Different weight vectors W transform the previous state and current input before combining

$$W^s \in \mathbb{R}^{H \times H}$$

$$W^x \in \mathbb{R}^{D \times H}$$

$$b \in \mathbb{R}^H$$

$$y_i = O(s_i) = s_i$$




Recurrent neural network

- Often used for sequential prediction tasks:
 - Language models—predicting the next symbol (word, character) in a sequence
 - Machine translation—predicting a sequence of words (sentence) in language f conditioned on sentence in language e
 - Sequence labeling (POS tagging, NER)

RNNs for sequence labeling

- The output state s_i is an H -dimensional real vector; we can transfer that into a probability by passing it through an additional linear transformation followed by a softmax

$$y_i = O(s_i) = \text{softmax}(s_i W^o + b^o)$$



Training RNNs

- Given this definition of an RNN:


$$s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$$

$$y_i = O(s_i) = \text{softmax}(s_iW^o + b^o)$$

- We have five sets of parameters to learn:

$$W^s, W^x, W^o, b, b^o$$

For POS tagging, predict the **tag** from ***y*** conditioned on the context



RNNs for POS

DT NN VBD IN DT NN ???

The horse raced past the barn fell

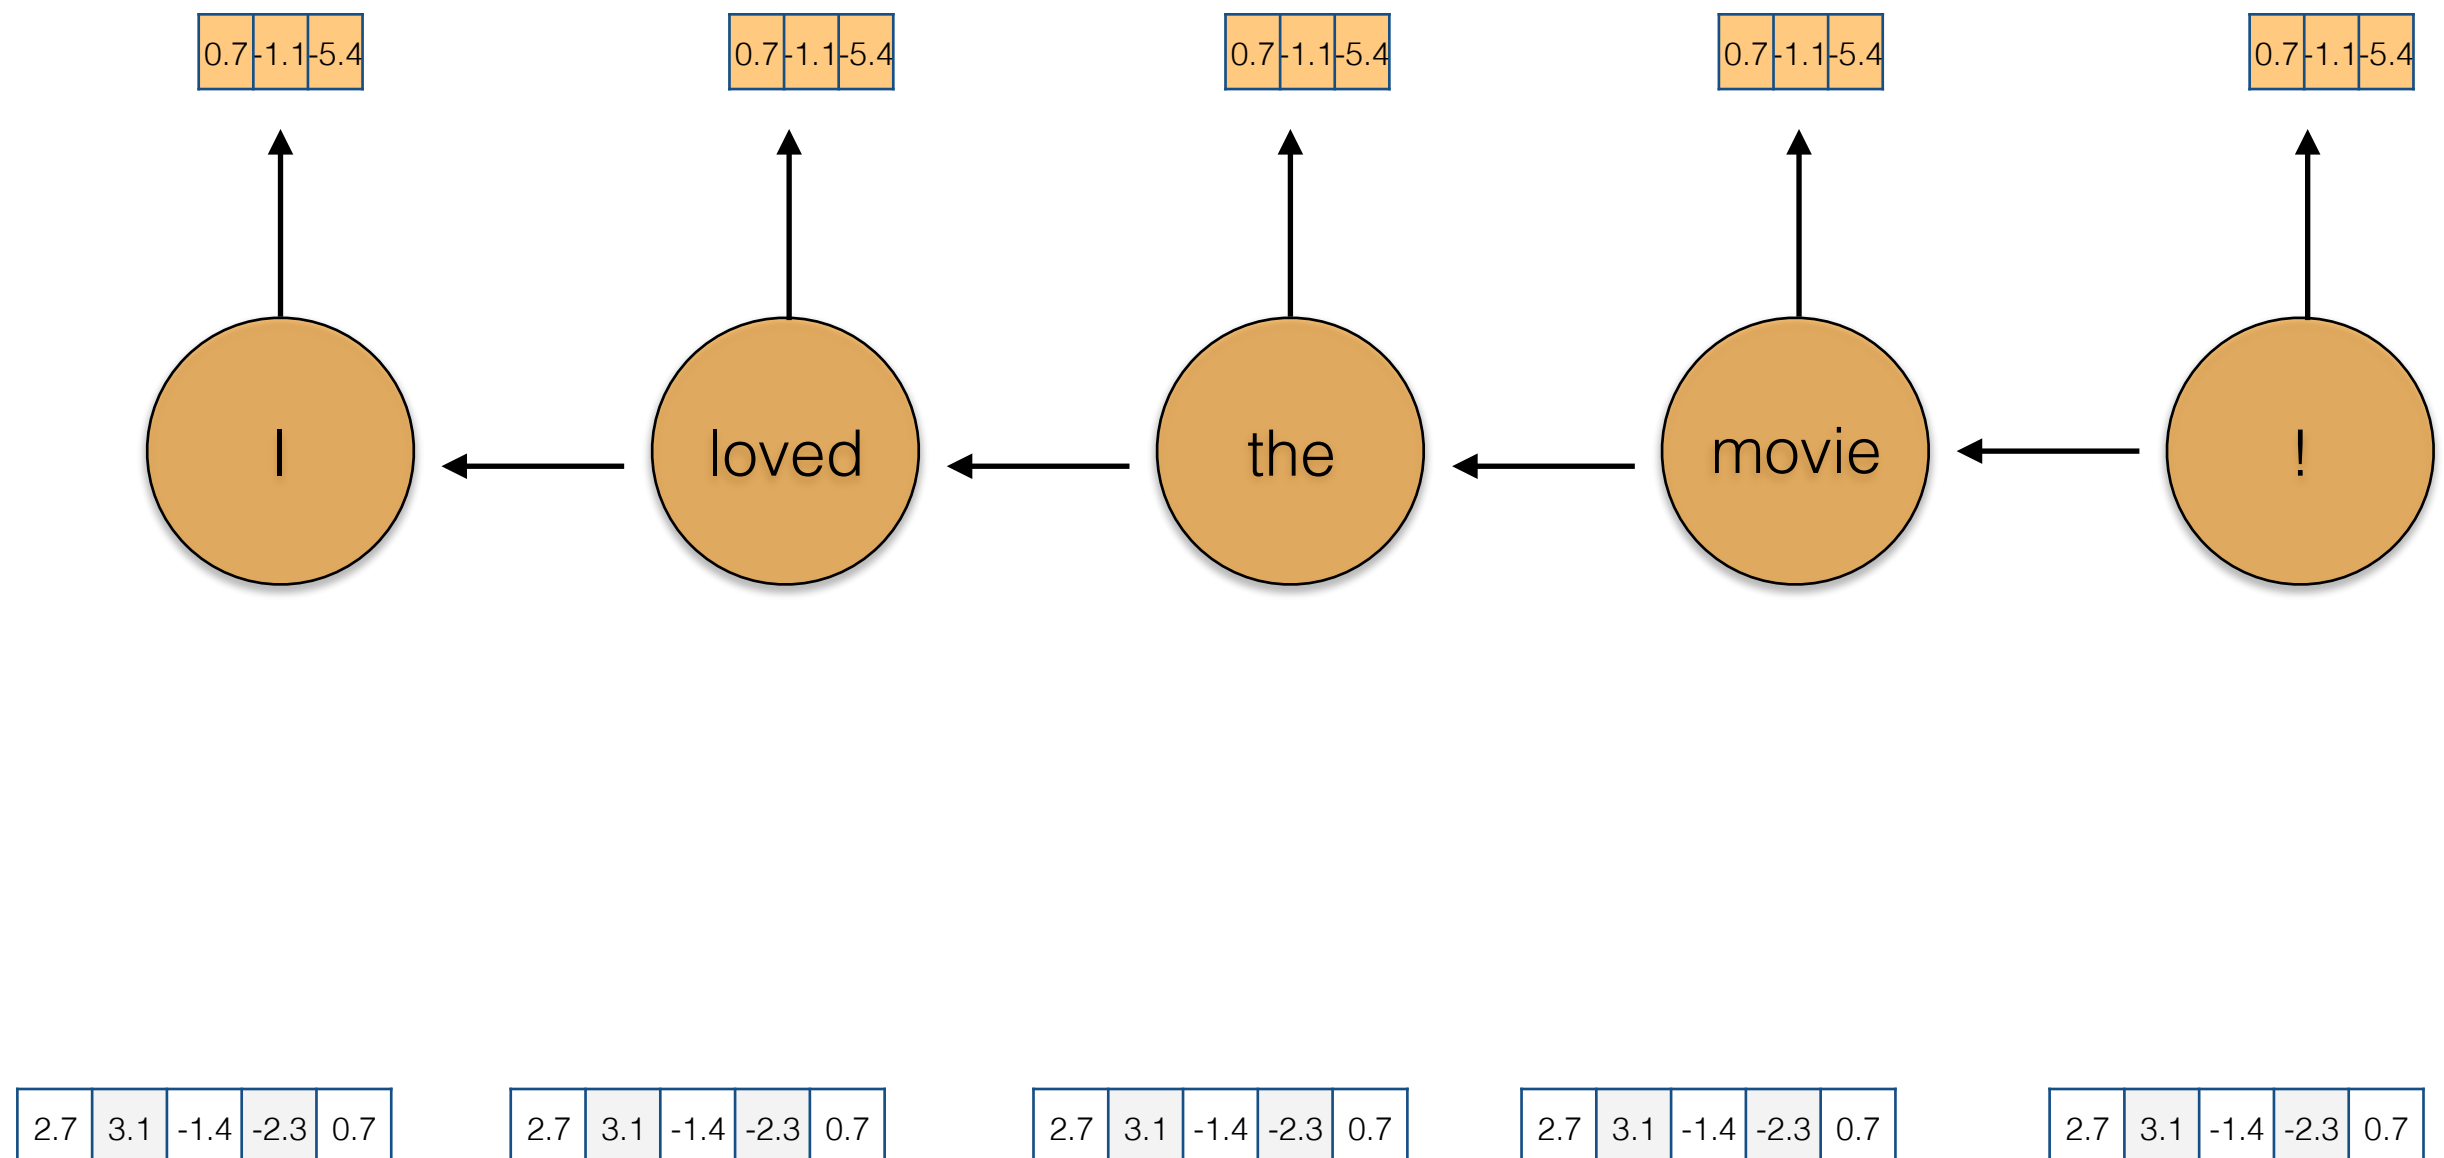
- To make a prediction for y_t , RNNs condition on all input seen through time t (x_1, \dots, x_t)
- But knowing something about the future can help (x_{t+1}, \dots, x_n)

Bidirectional RNN

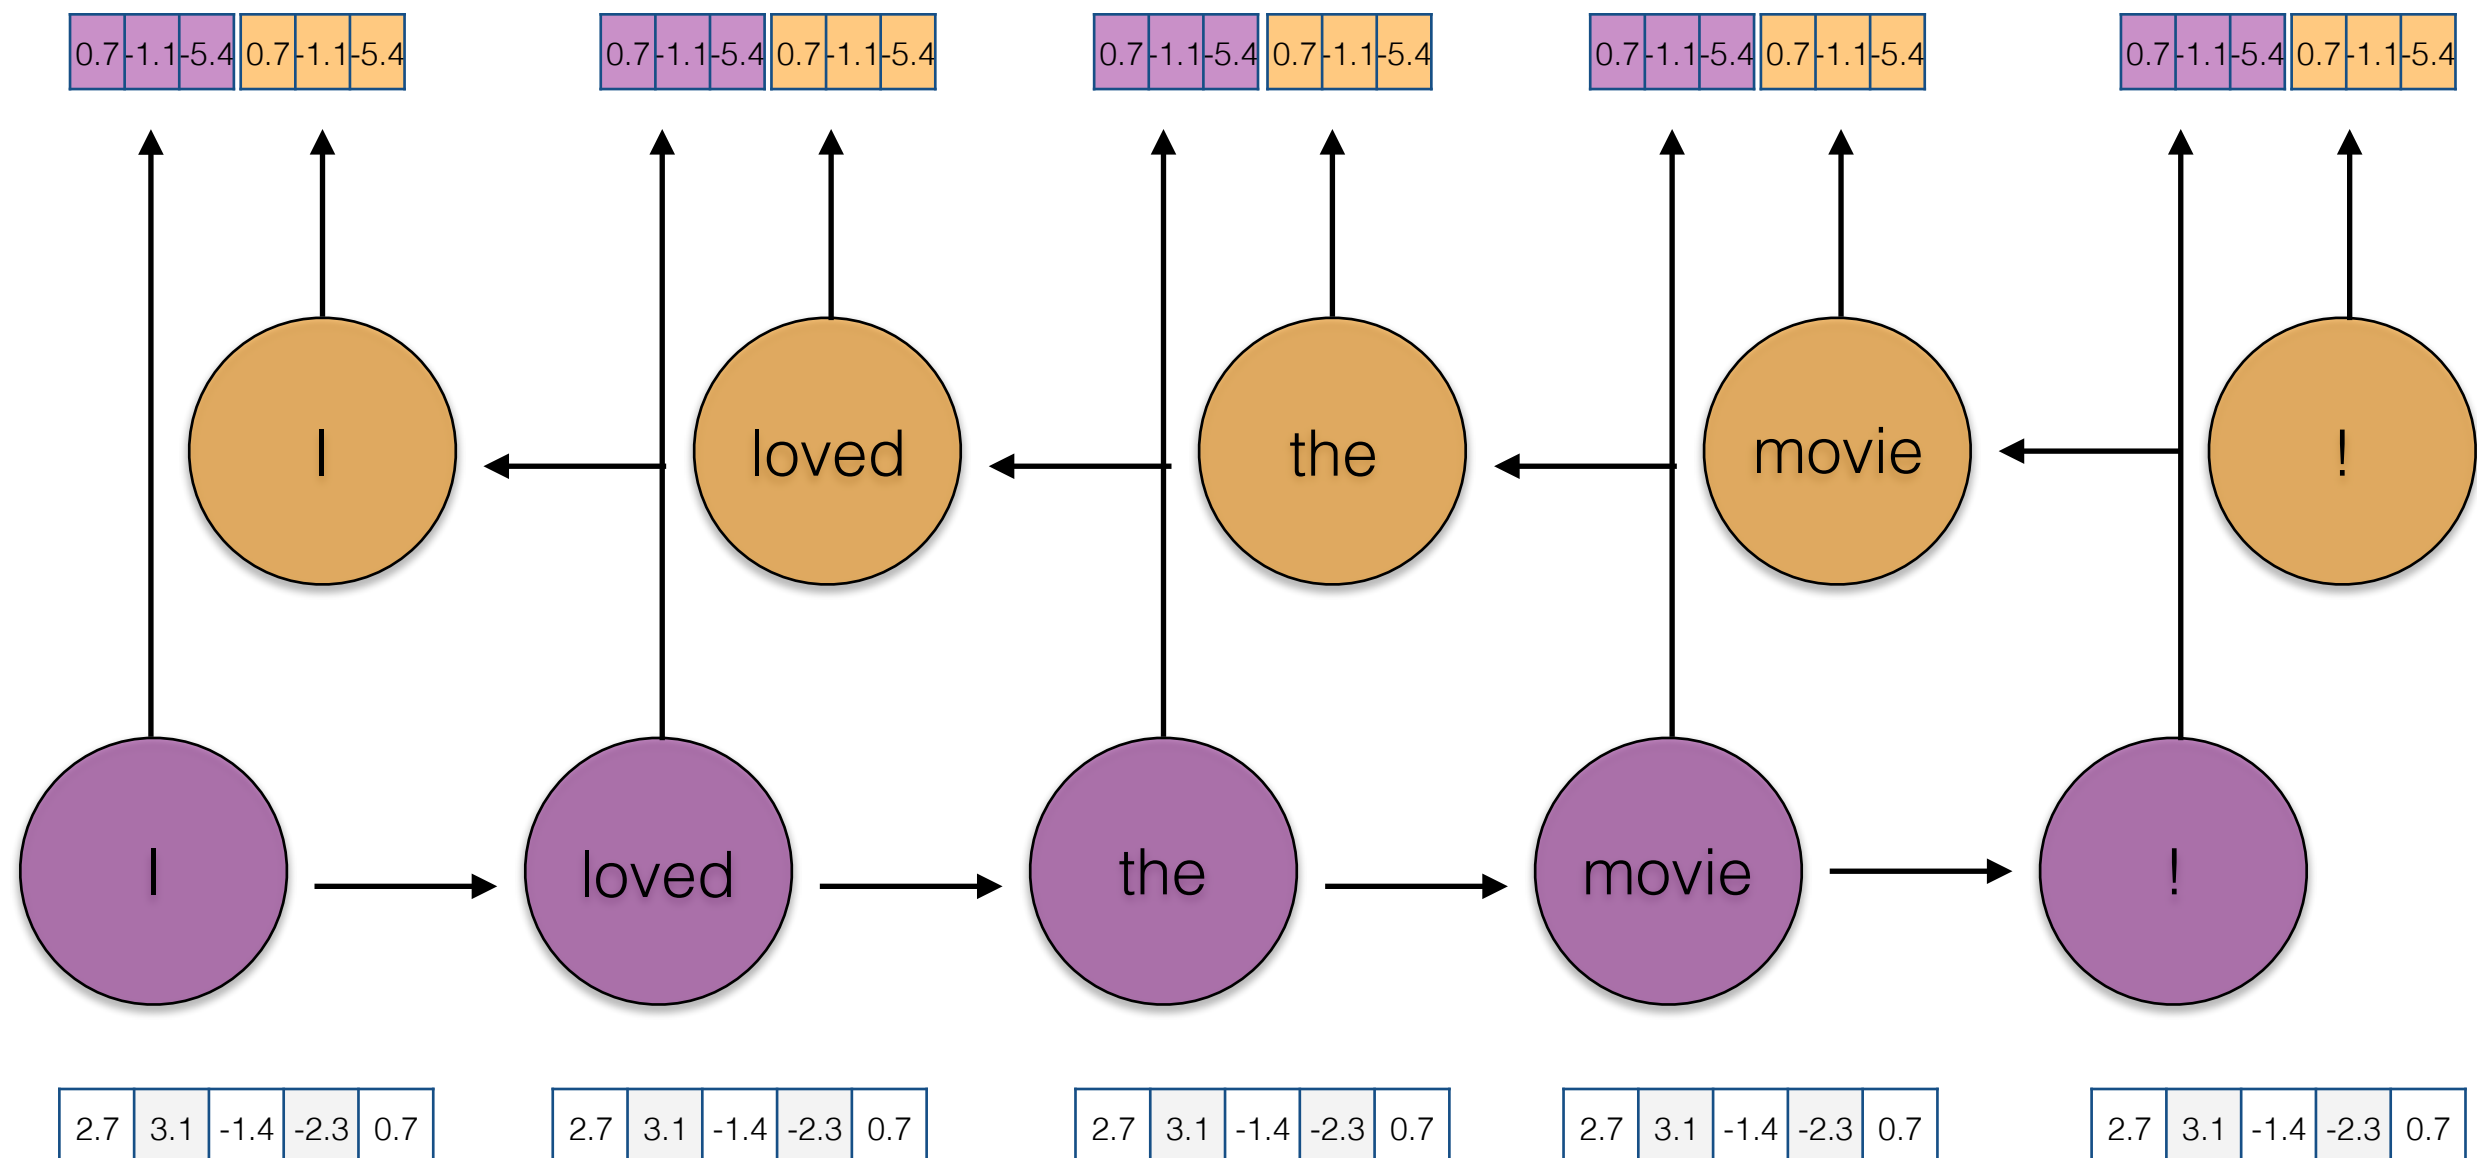
- A powerful alternative is make predictions conditioning both on the **past** and the **future**.
- Two RNNs
 - One running left-to-right
 - One right-to-left
- Each produces an output vector at each time step, which we concatenate

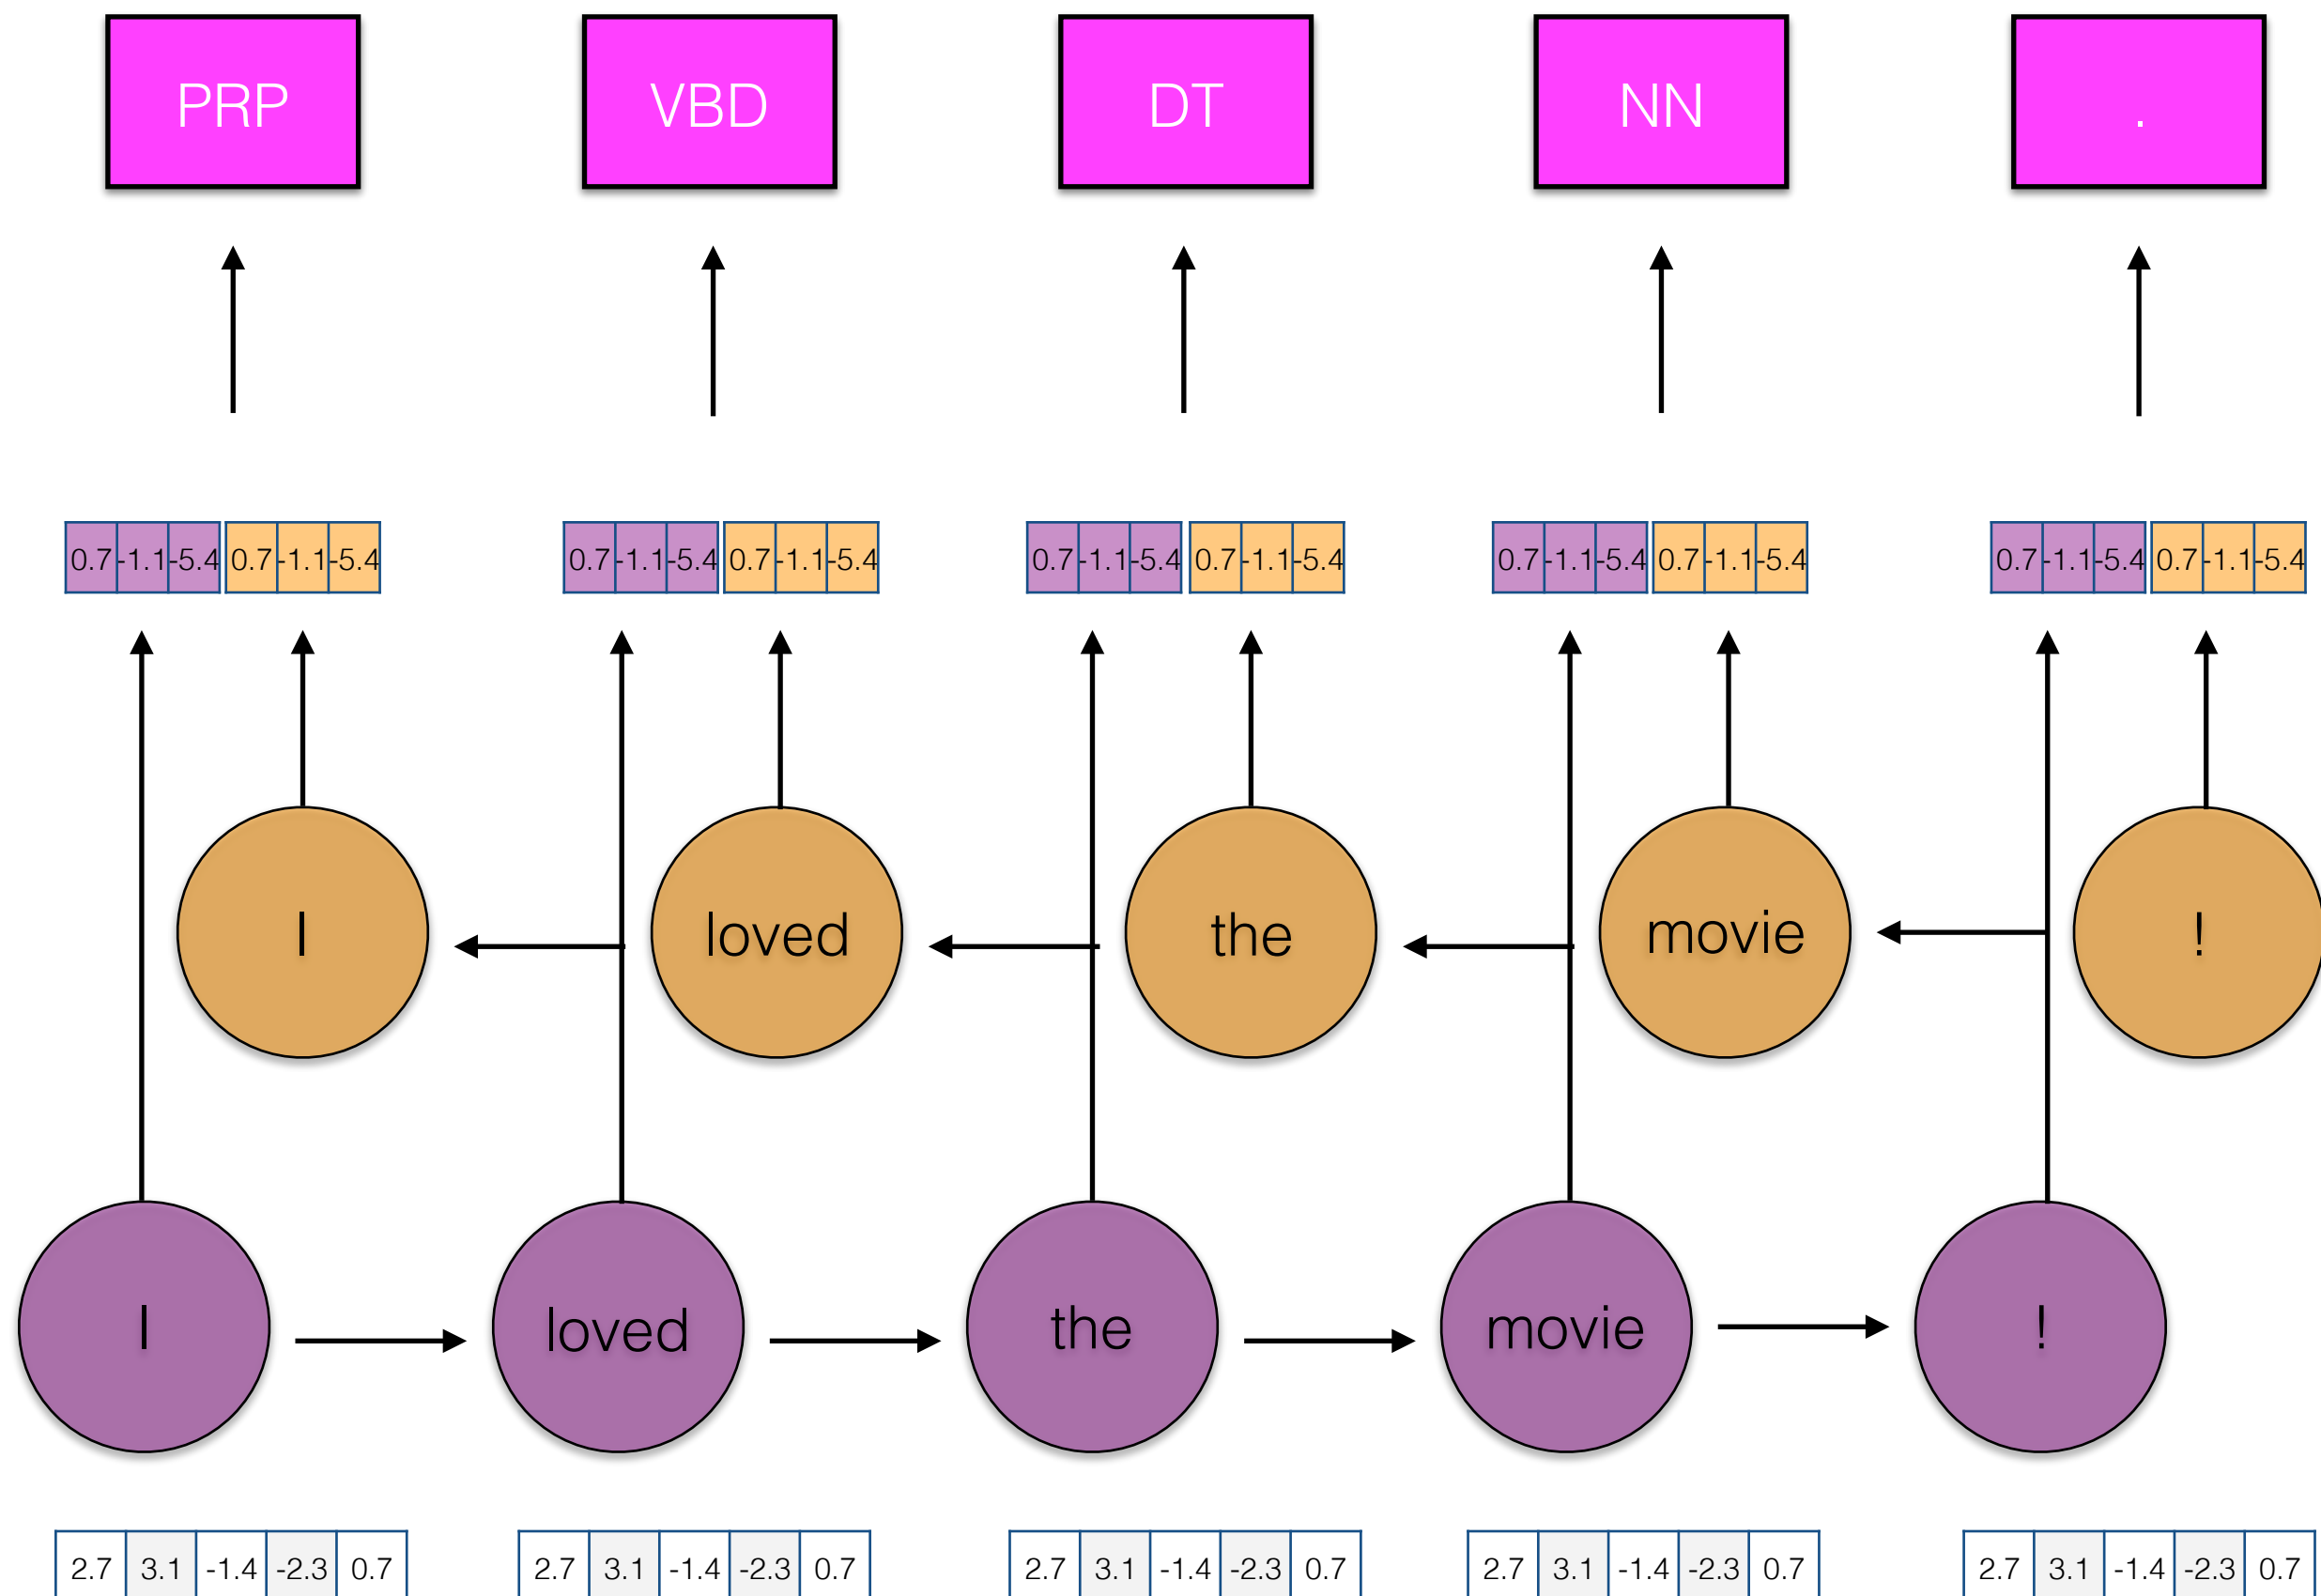
Bidirectional RNN

backward RNN



Bidirectional RNN





Training BiRNNs

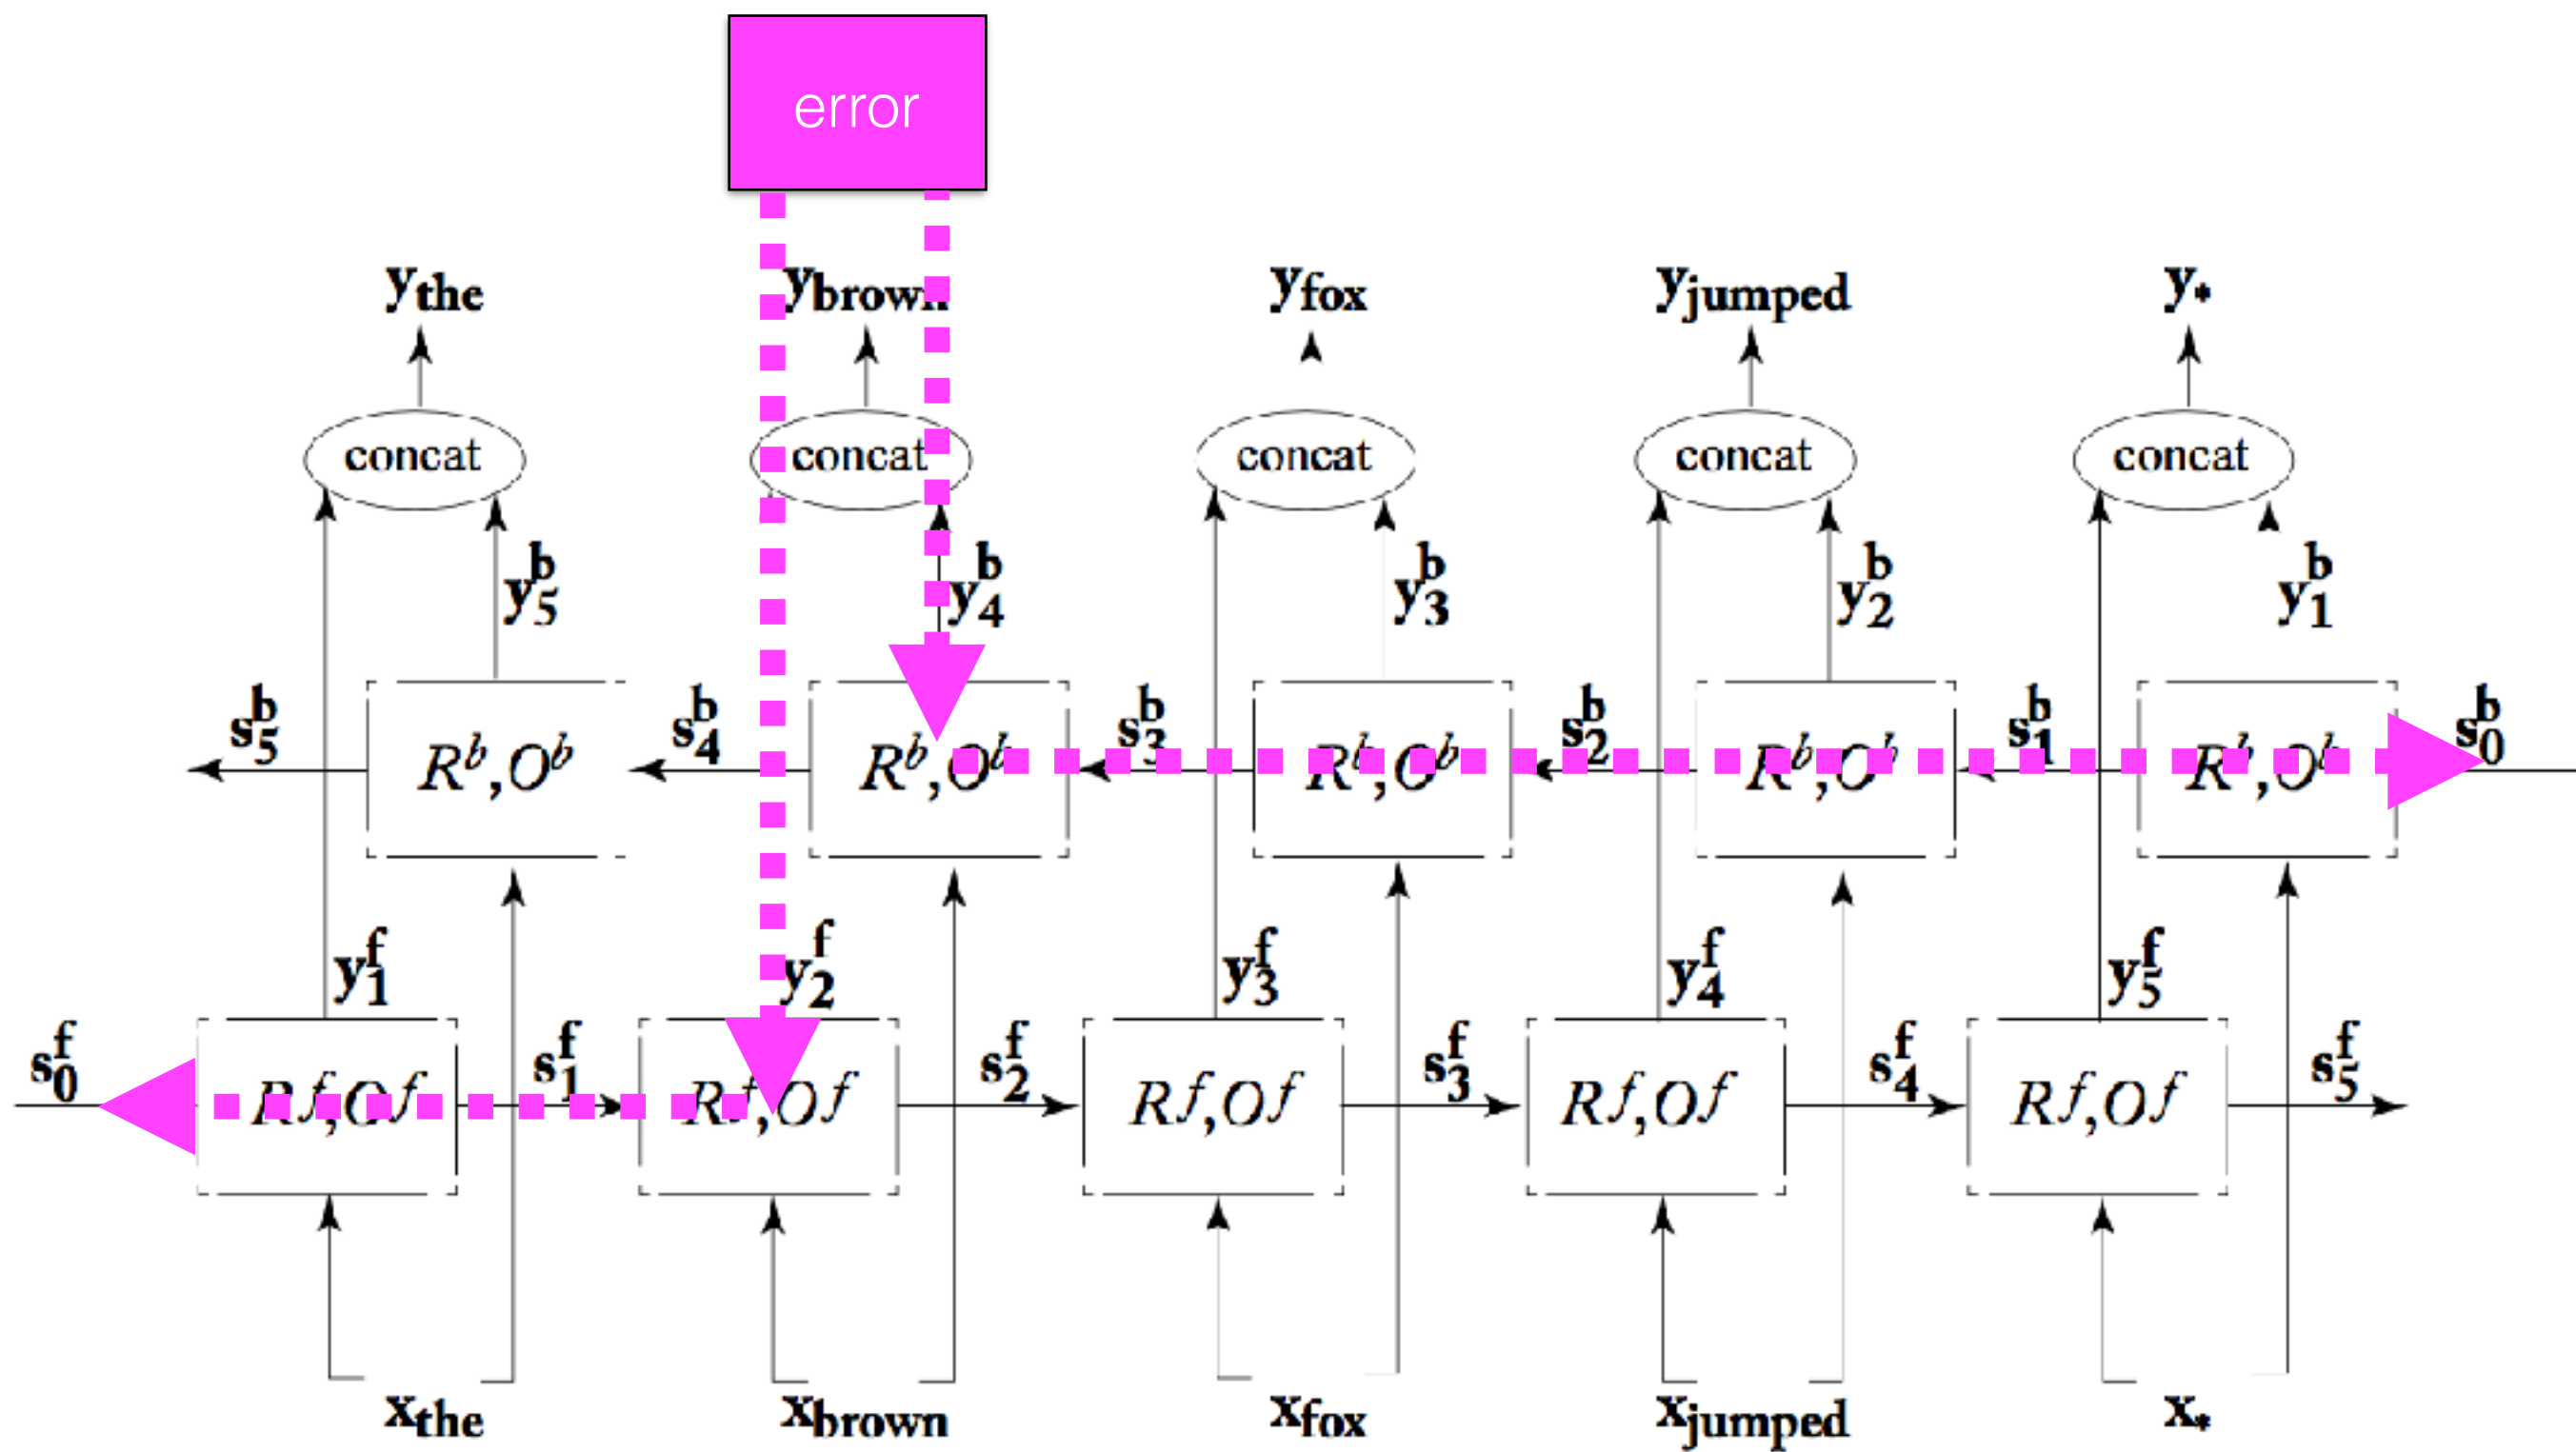
- Given this definition of an BiRNN:

$$s_b^i = R_b(x^i, s_b^{i+1}) = g(s_b^{i+1} W_b^s + x^i W_b^x + b_b)$$

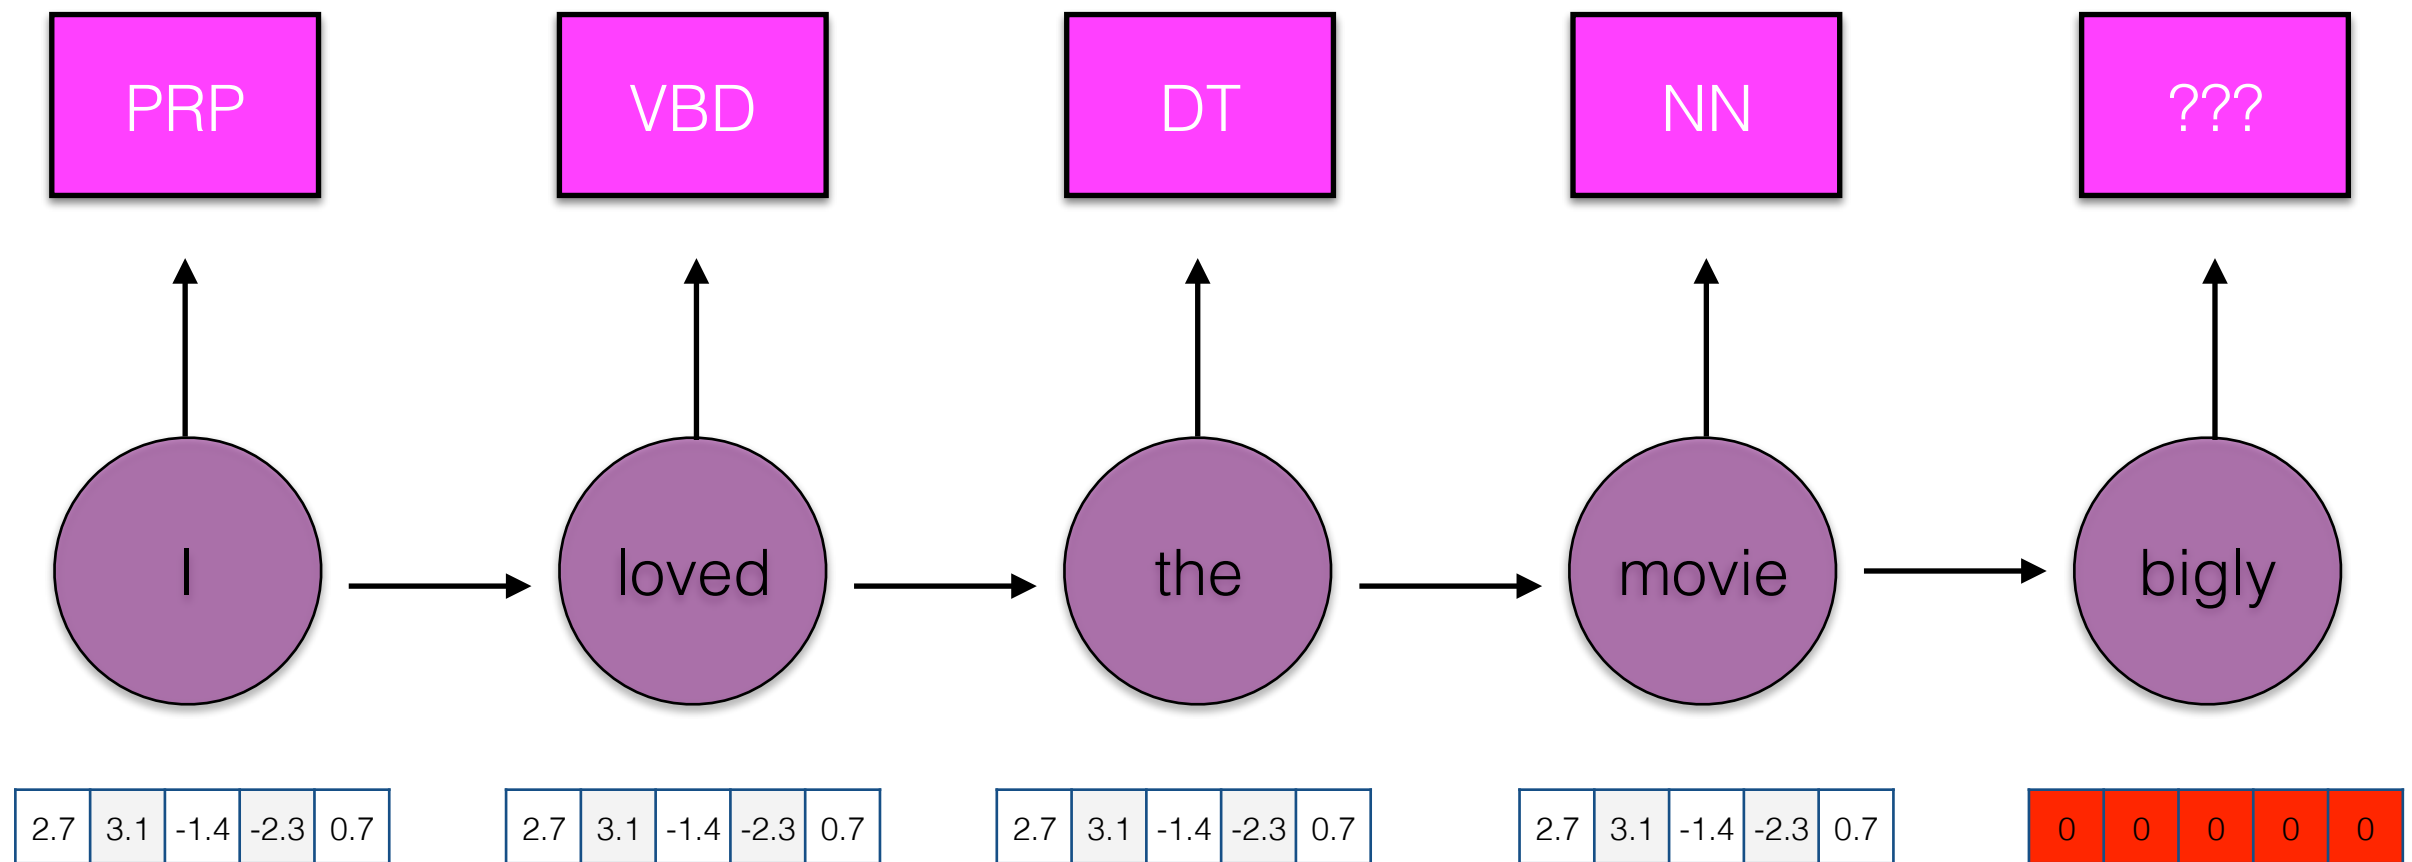
$$s_f^i = R_f(x^i, s_f^{i-1}) = g(s_f^{i-1} W_f^s + x^i W_f^x + b_f)$$

$$y_i = \text{softmax}([s_f^i; s_b^i] W^o + b^o)$$

- We have 8 sets of parameters to learn (3 for each RNN + 2 for the final layer)



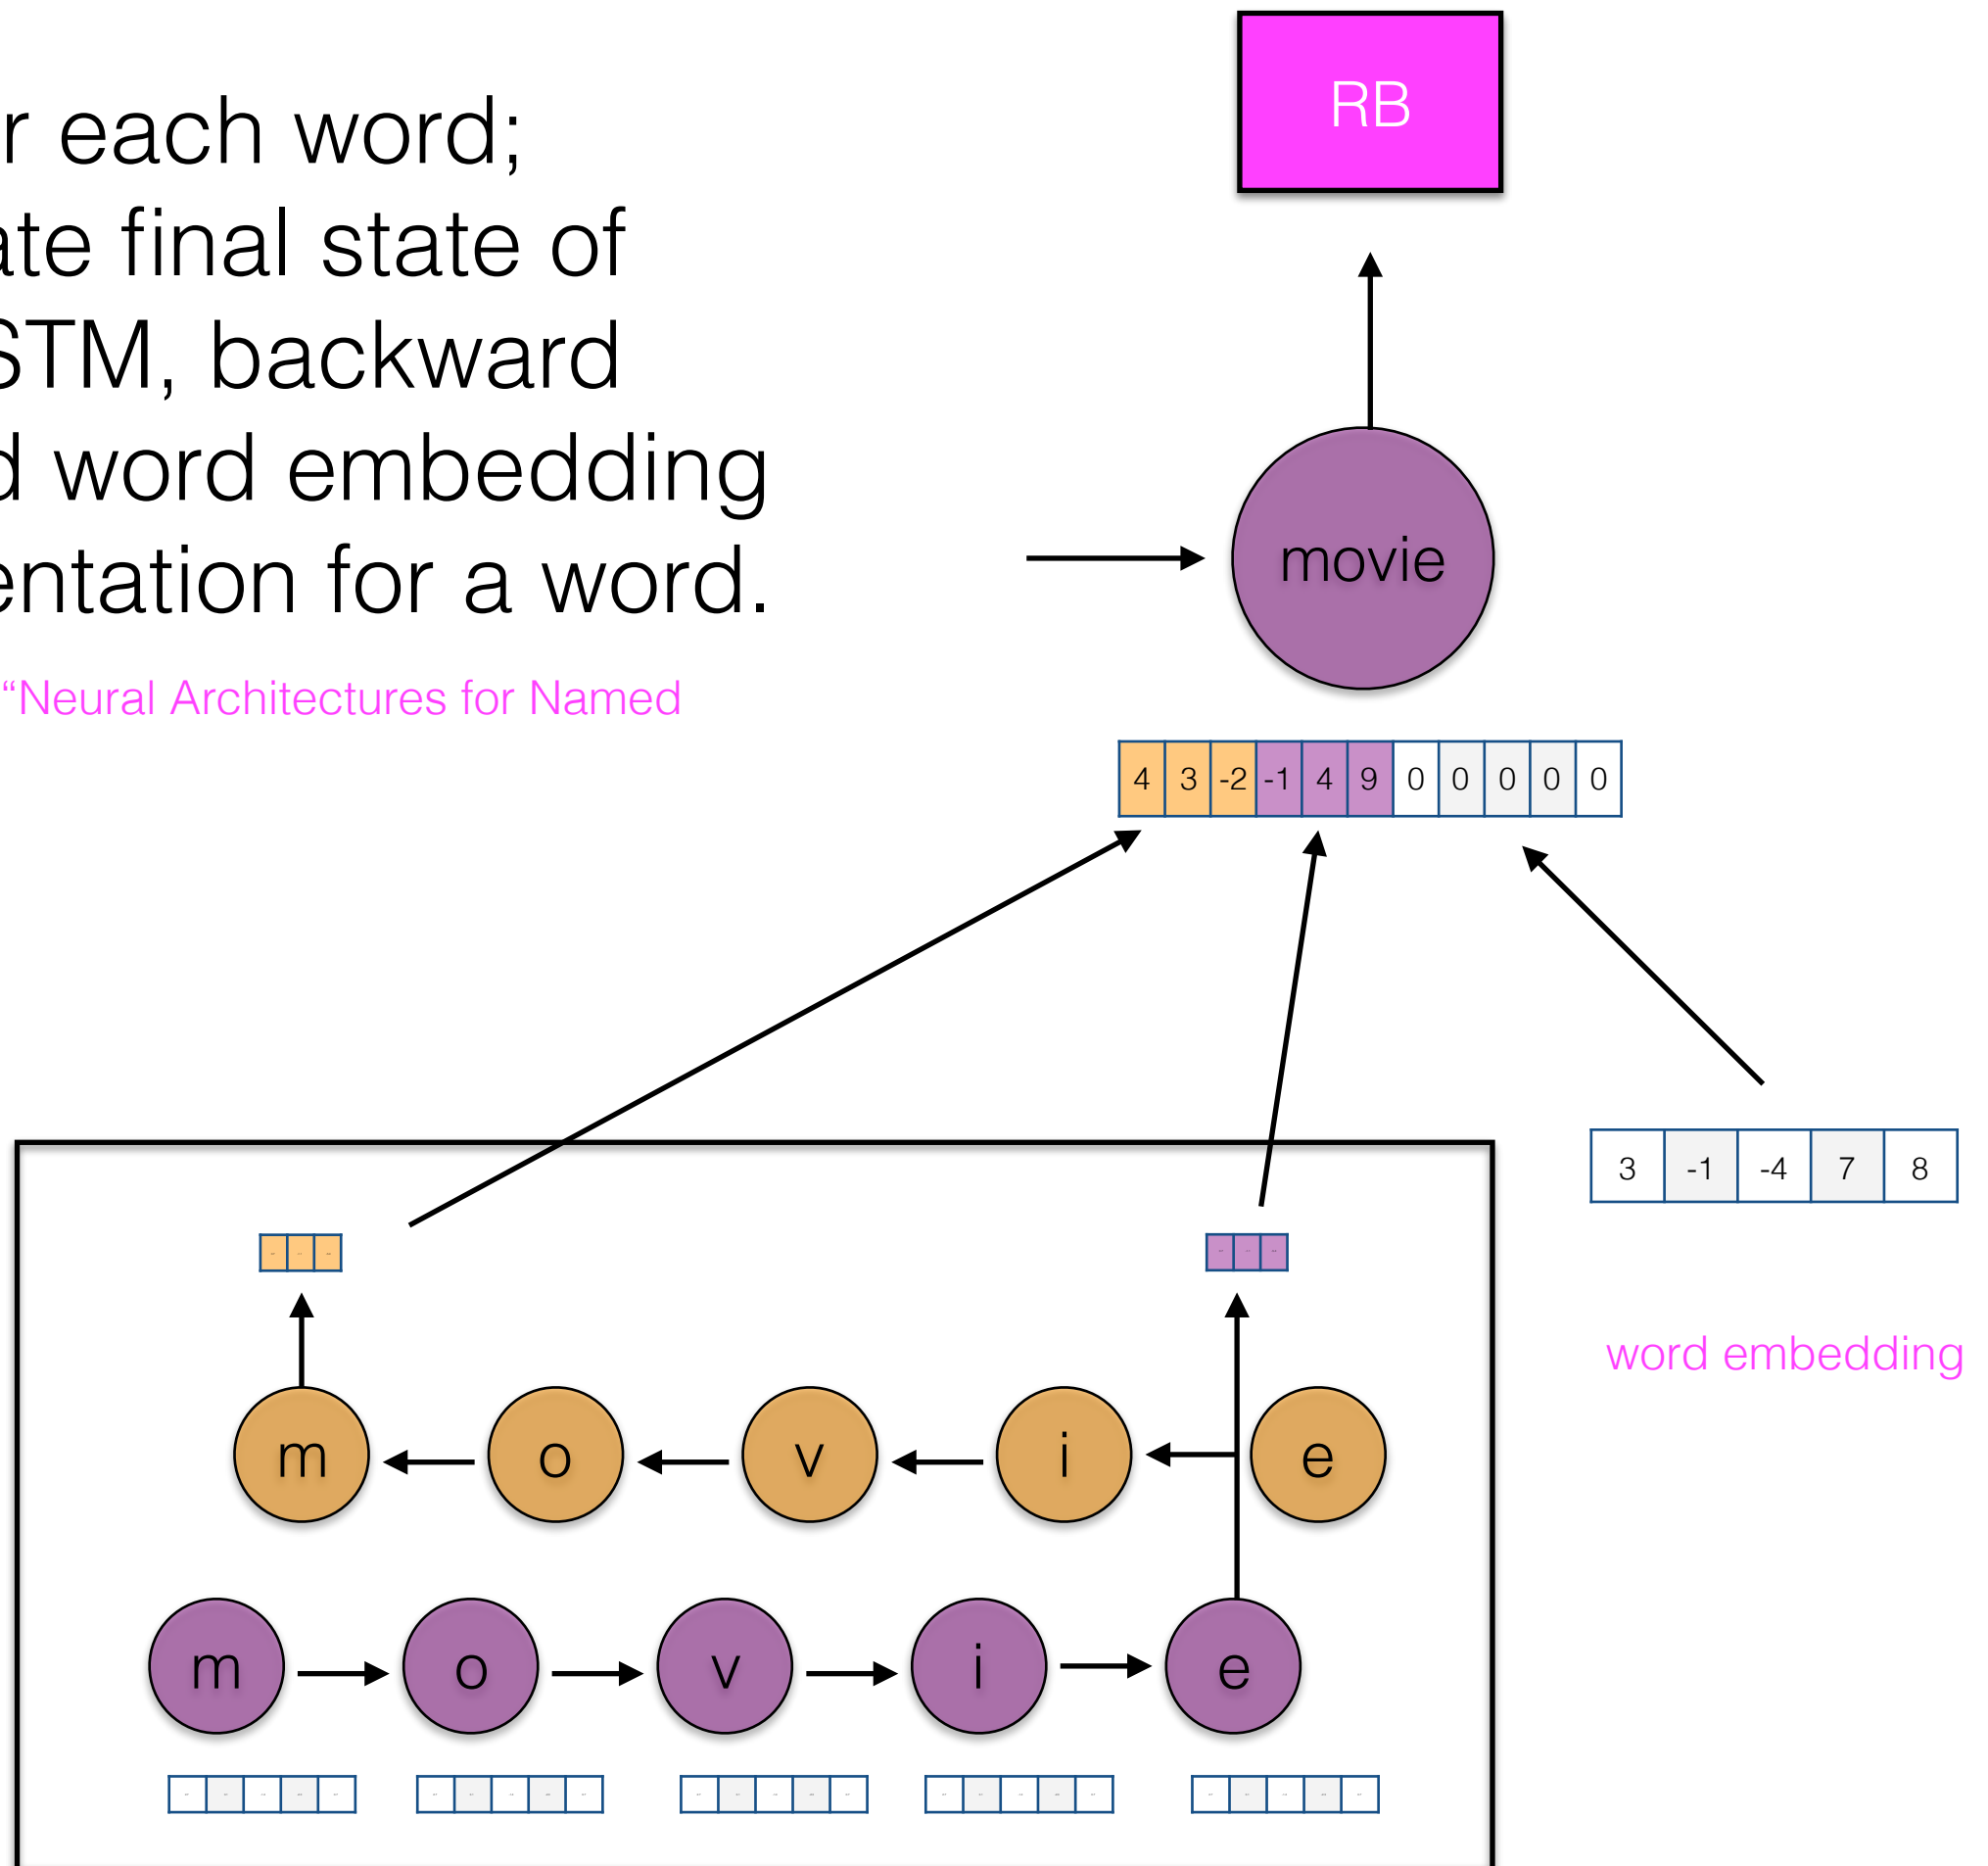
How do we fix this?



BiLSTM for each word;
concatenate final state of
forward LSTM, backward
LSTM, and word embedding
as representation for a word.

Lample et al. (2016), "Neural Architectures for Named
Entity Recognition"

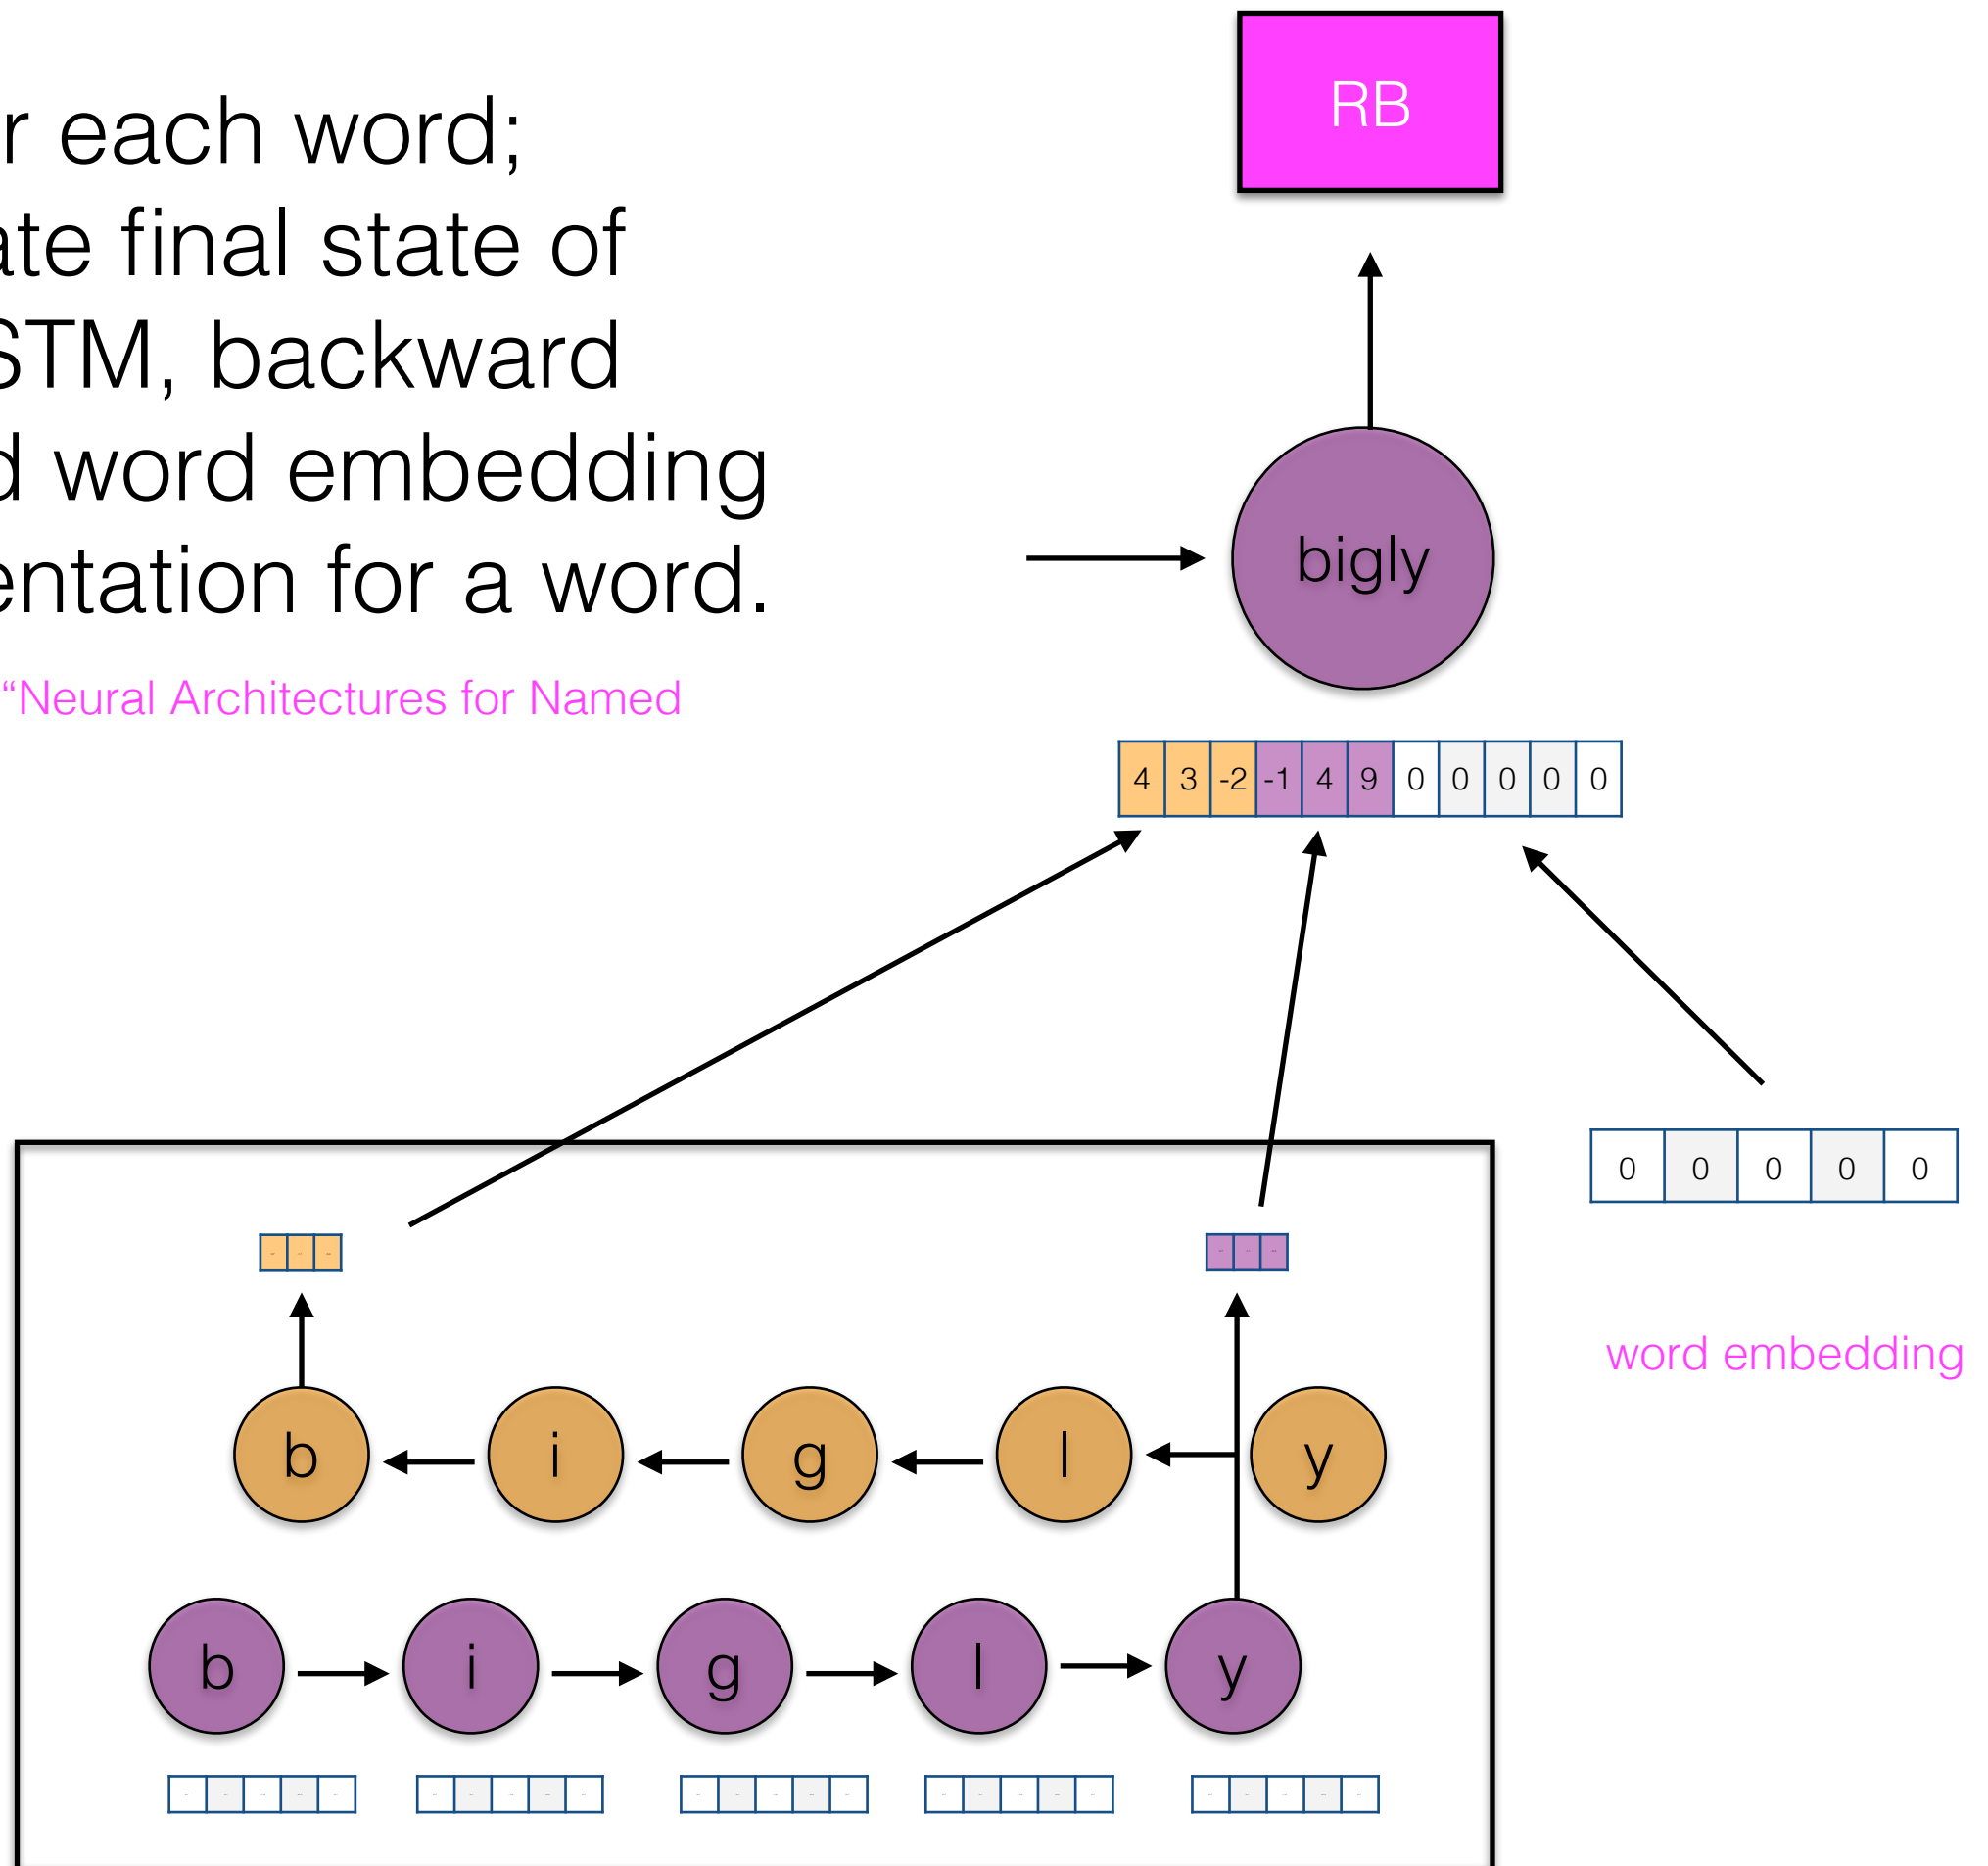
character BiLSTM



BiLSTM for each word;
concatenate final state of
forward LSTM, backward
LSTM, and word embedding
as representation for a word.

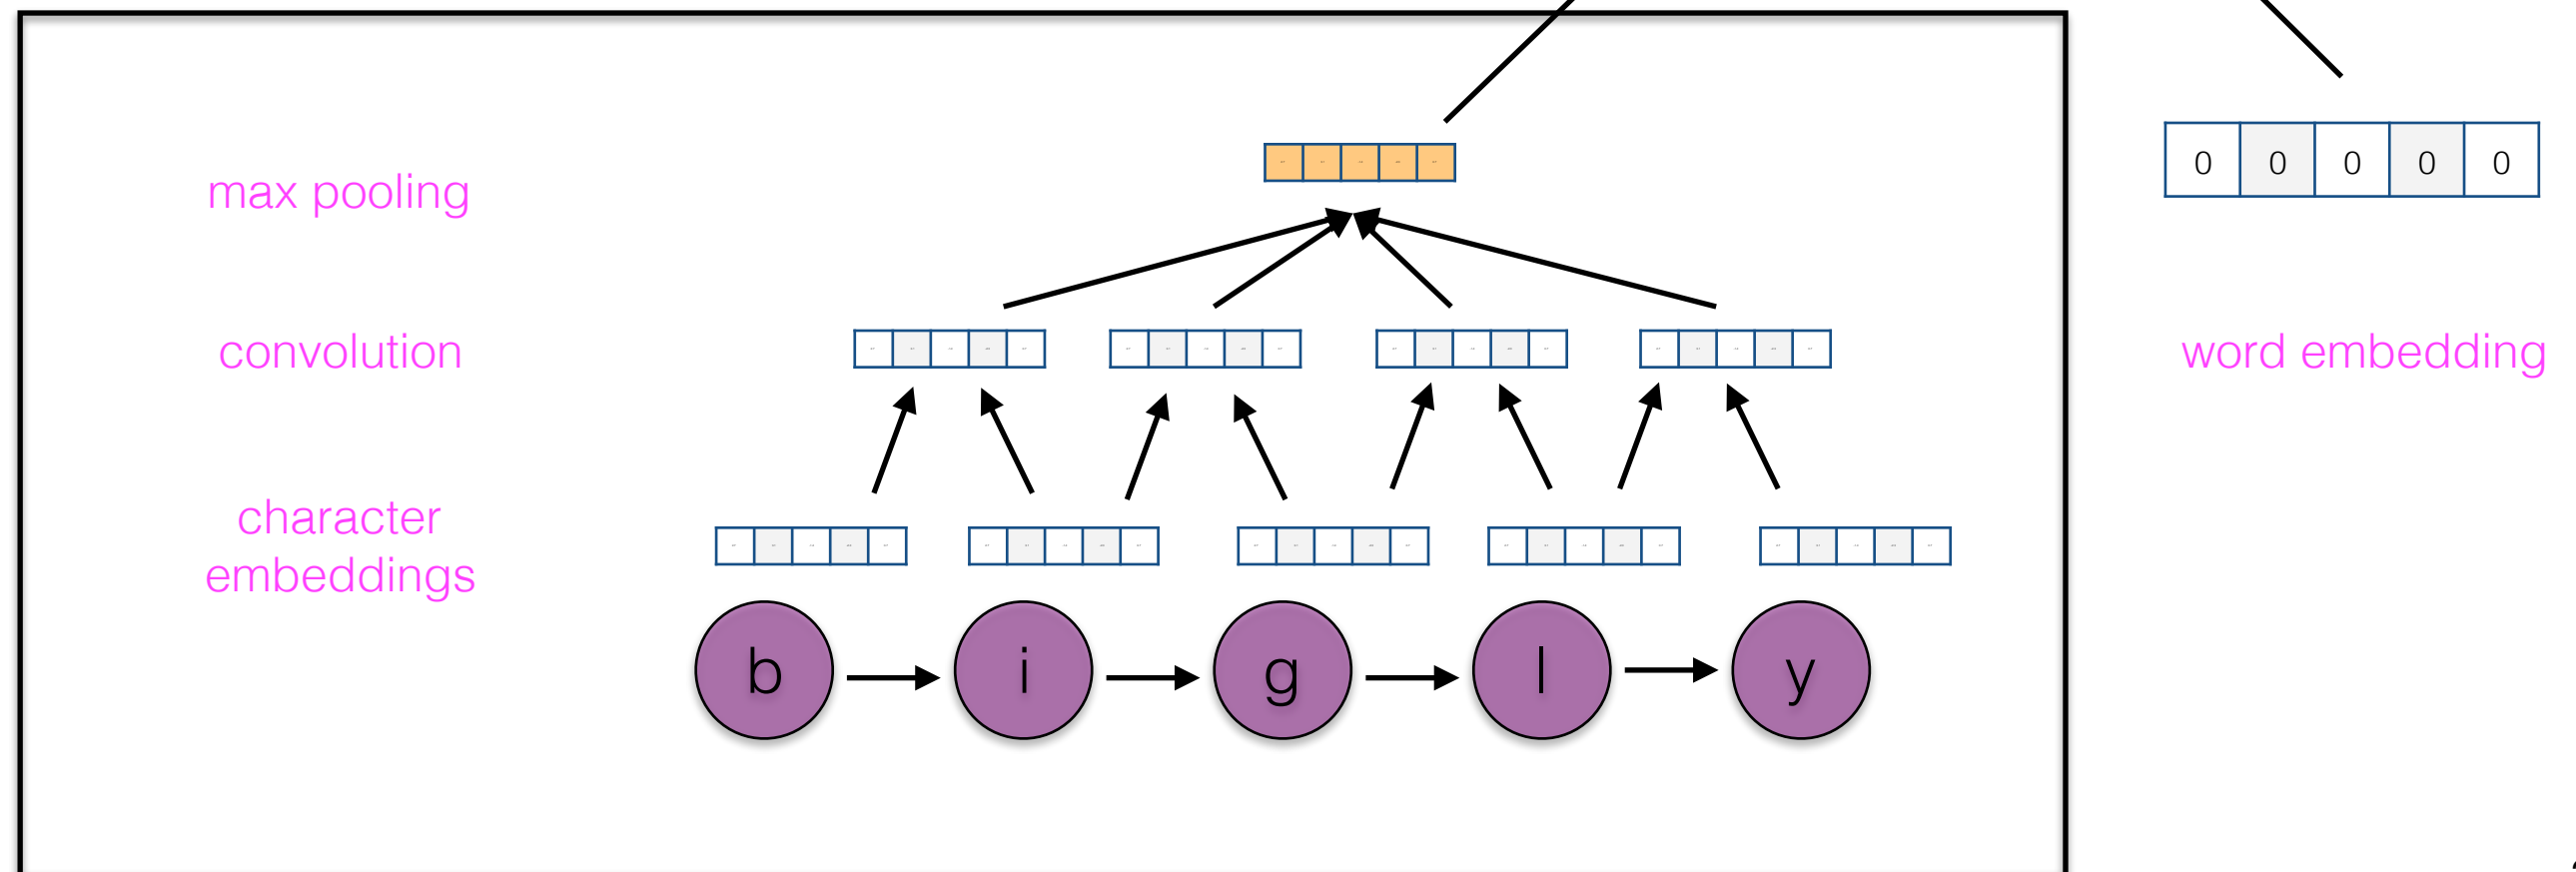
Lample et al. (2016), "Neural Architectures for Named
Entity Recognition"

character BiLSTM



Character CNN for each word;
concatenate character CNN
output and word embedding
as representation for a word.

Chu et al. (2016), "Named Entity Recognition with
Bidirectional LSTM-CNNs"



LSTM/RNN

- An RNN doesn't use the dependencies between nearby **labels** in making predictions.


NN

VBZ

VB

VBZ	0.51
NNS	0.48
JJ	0.01
NN	0
...	...

The information that's passed between states is not the categorical choice (VBZ) but a hidden state that generated the distribution.



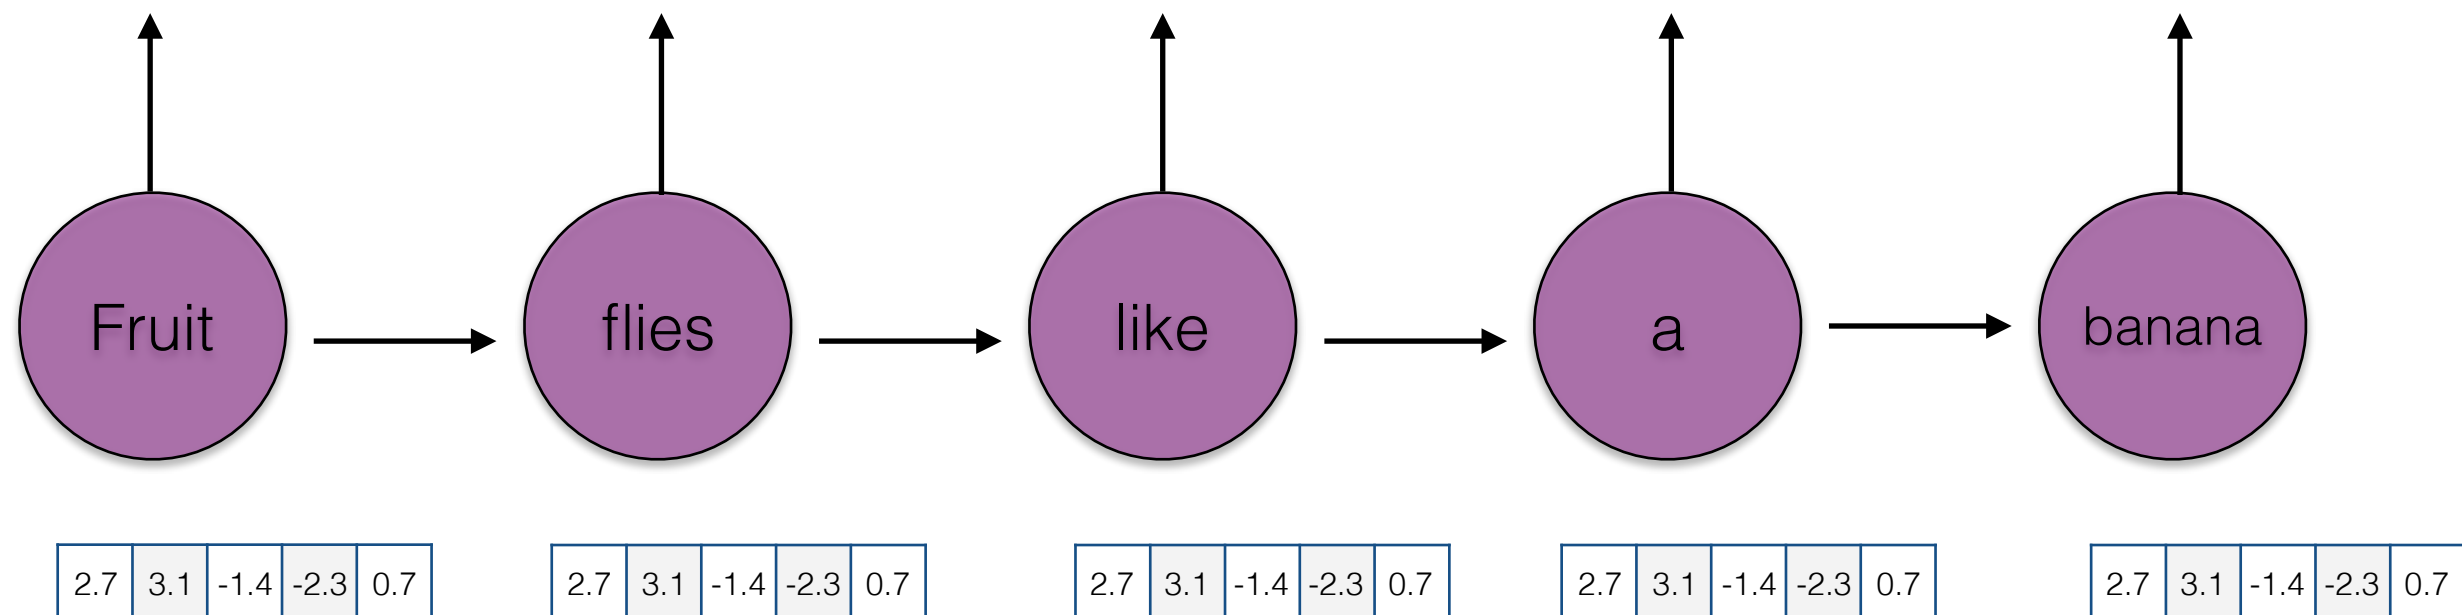
NN

VBZ

VB

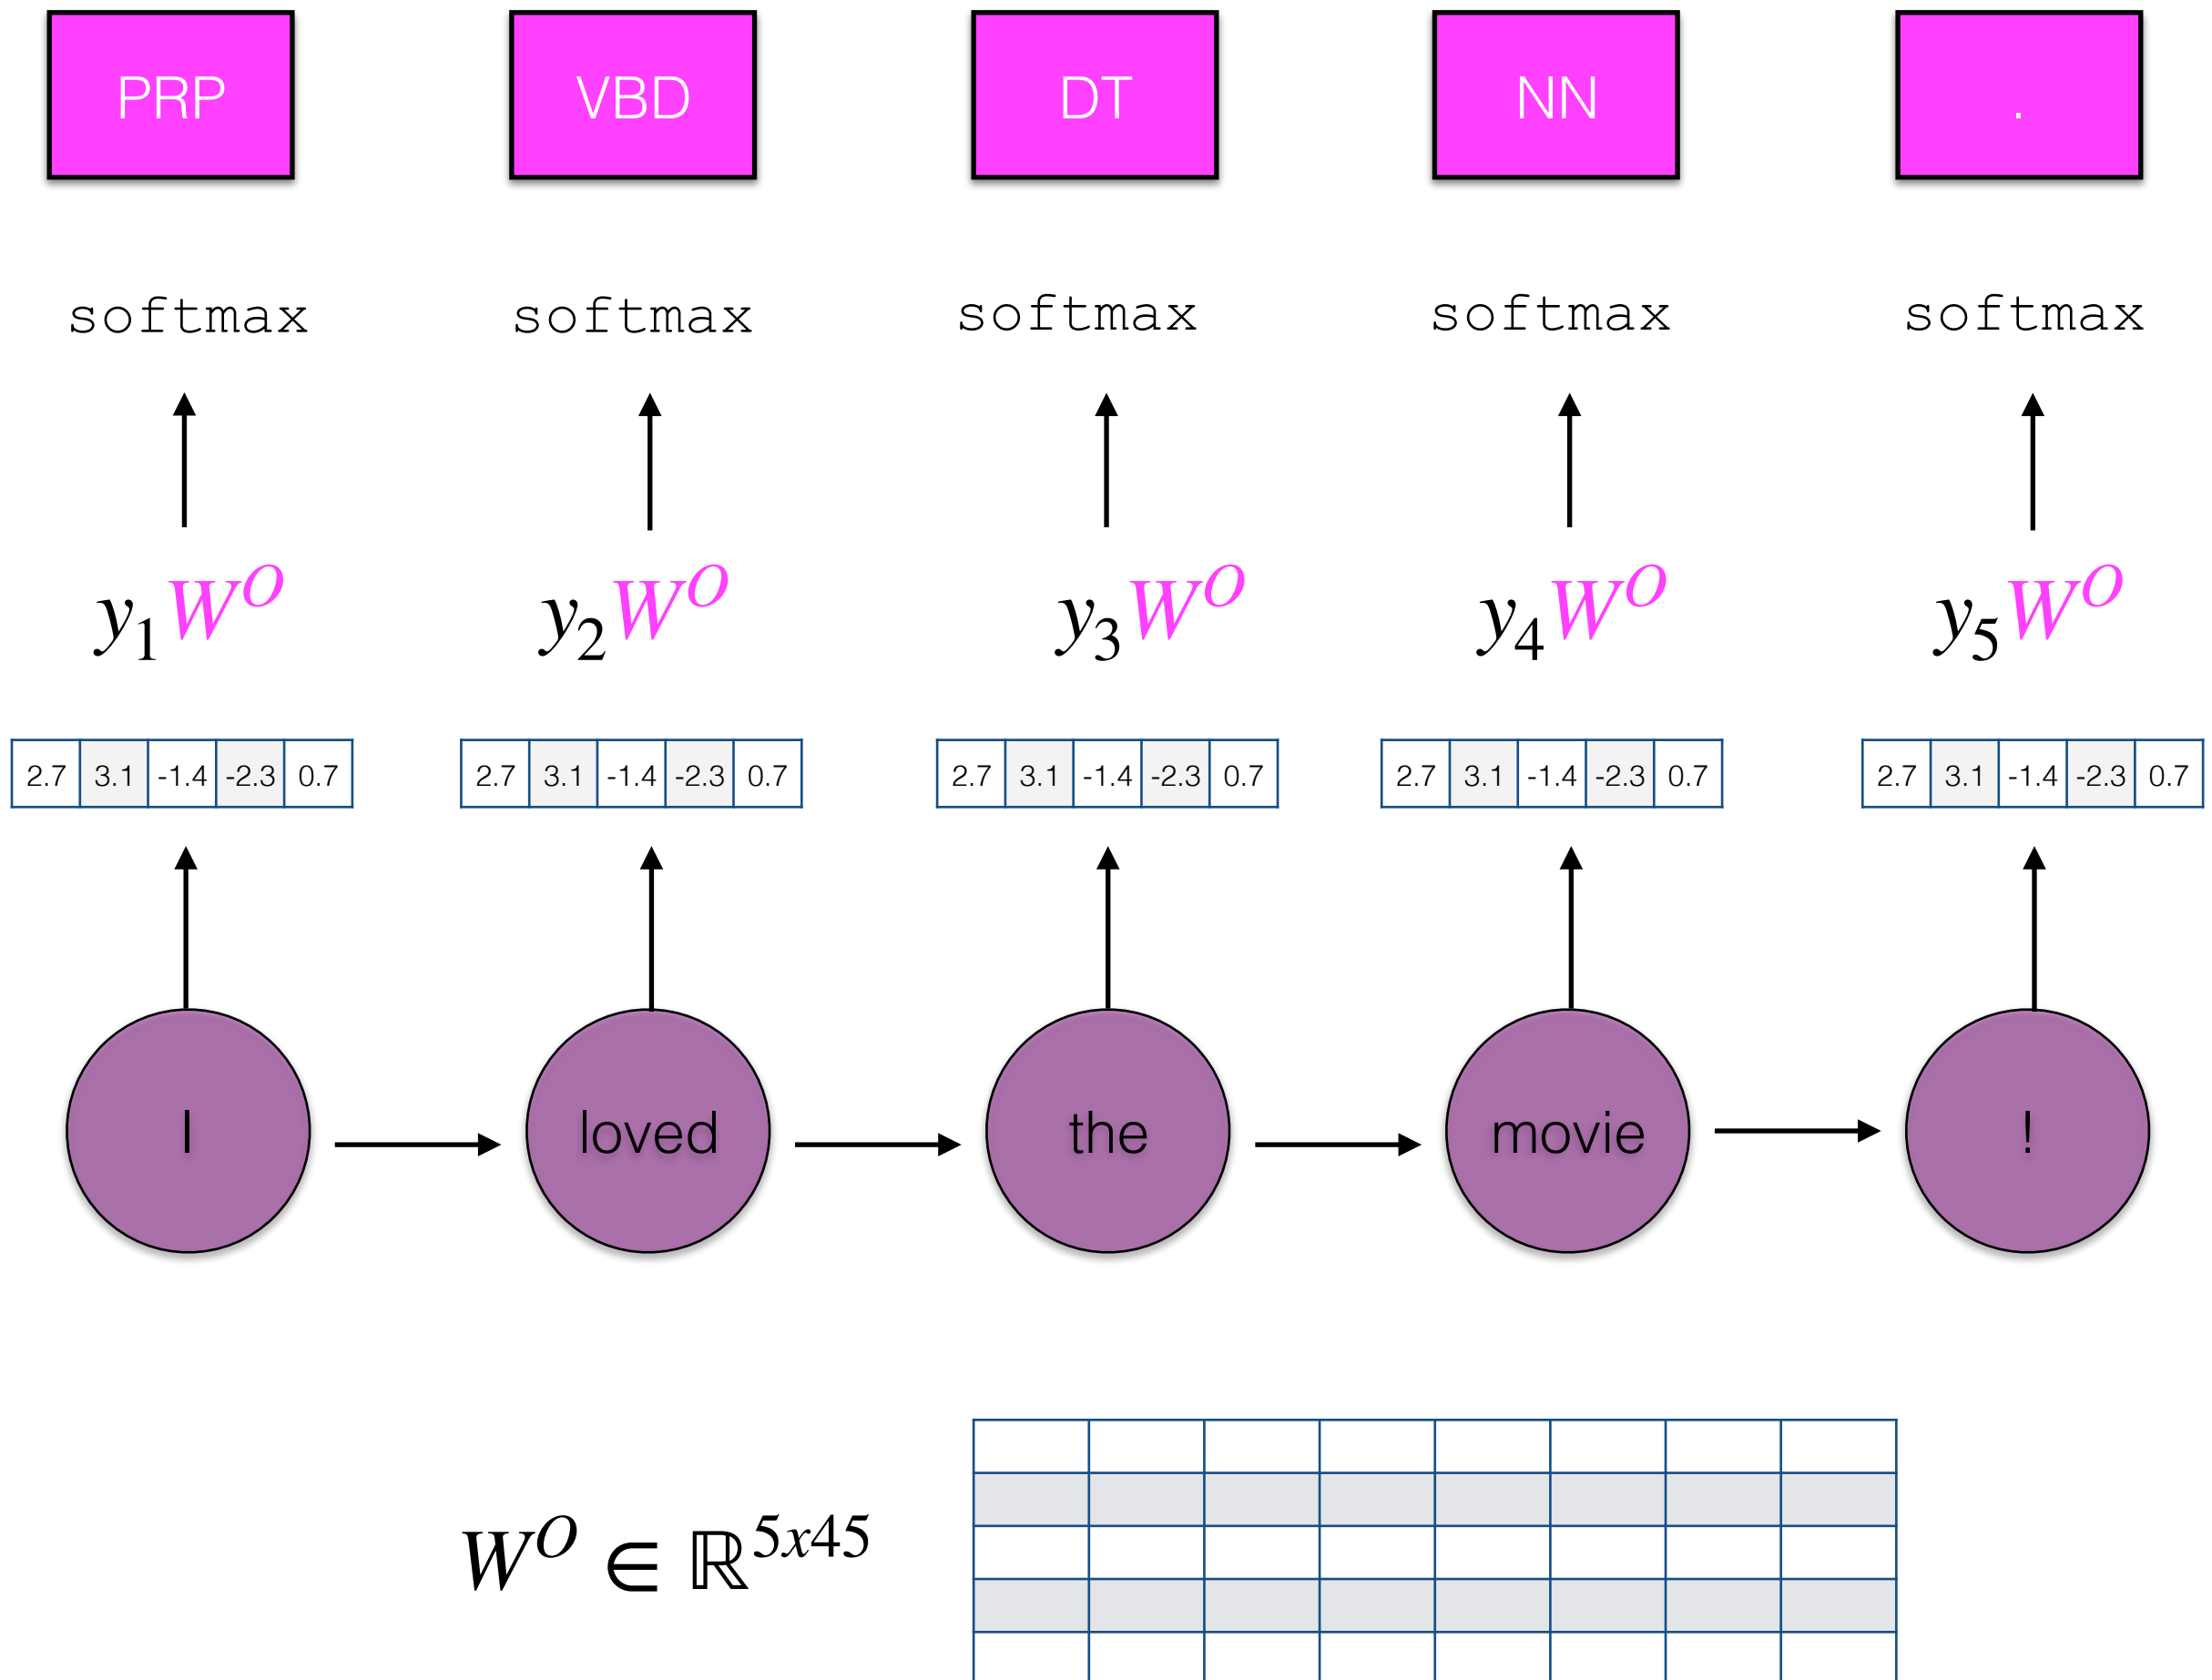
VBZ	0.51
NNS	0.48
JJ	0.01
NN	0
...	...

If we knew the categorical choice of **VBZ** at t_2 , $P(\text{VB})$ at t_3 would be much lower.



TimeDistributed

- In keras, the `TimeDistributed` wrapper applies the same operation to every time step in a sequence (e.g., the same `Dense` layer with the same parameters)



TimeDistributed

```
lstm_output = LSTM(lstm_size,  
return_sequences=True)  
(embedded_sequences)
```

```
preds = TimeDistributed(Dense(output_dim,  
activation="softmax"))(lstm_output)
```

Activity

- 12.ner/SequenceLabelingBiLSTM_TODO